

ATFD: Agent Trajectory Failure Detection

A Benchmark for Evaluating Agent Monitoring Tools

Anonymous Authors

anonymous@example.com

Abstract

Agent workflows built on large language models are increasingly deployed in business-critical settings, yet no benchmark exists for evaluating the *monitoring tools* that must detect failures in these workflows. Existing benchmarks evaluate agent *capability*—whether an agent completes a task—but not whether an external system can reliably *detect* when an agent has failed, partially succeeded, or produced substandard output. We introduce ATFD (Agent Trajectory Failure Detection), a benchmark that fills this gap. ATFD contributes: (1) a 7-category, 23-subcategory failure taxonomy that extends prior work by including quality degradation as a first-class failure mode; (2) a formal task definition where systems must analyze complete agent trajectories and produce structured findings with severity, category, and attribution; (3) a multi-domain dataset of 2,462 trajectories spanning τ -bench (customer service), SWE-BENCH (coding), hand-crafted synthetic scenarios, ATBench (agent safety, 200-sample subset), and Toolathlon (long-horizon tool use); (4) a multi-source ground truth protocol combining programmatic verdicts with LLM-judge consensus; (5) a complete comparative evaluation of 7 systems across all 6 domains—a naive heuristic baseline, a zero-config structural heuristic (Galea), two open-weight LLM judges (Llama 4 Scout and Llama 3.3 70B), Claude Sonnet 4 as a proprietary LLM judge, LangSmith with custom evaluator rules, and Braintrust with custom scorers—revealing a stark configuration effort gap and clear domain difficulty spectrum; and (6) an open benchmark harness with cost-aware metrics. Our complete evaluation reveals a stark *detection gap* with complementary profiles: the Galea heuristic achieves 100% detection rate on all τ -bench domains with 0% false positives using zero-config pattern matching, but only 78% on synthetic failures that require semantic analysis; conversely, open-weight LLM judges achieve 100% detection on synthetic failures but struggle on real τ -bench failures (Llama 4 Scout: 25.0% on retail, 20.0% on airline). ATBench (200-sample subset of 1,000 safety trajectories) reveals a middle ground for Galea (52.5% DR between its structural strengths and limits) and full 100% DR for Claude Sonnet 4; Galea reaches 52.5% by detecting structurally co-occurring anomalies, while the remaining 47.5% require semantic understanding. Toolathlon (100 long-horizon trajectories with 600+ tools, up to 16k messages) is the hardest domain for *all* systems except Claude Sonnet 4, which uniquely achieves 100% DR even here, though at a high FPR of 67%; every other evaluated system achieves 0% DR, confirming that long-horizon multi-step failures require frontier semantic understanding and cannot be caught by structural analysis or rule-based approaches. SWE-BENCH emerges as the “easy” domain: its failures are structurally obvious (tool errors, test failures) and detectable at 100% DR even by the naive baseline, contrasting sharply with airline’s semantically subtle failures where even Claude Sonnet 4 reaches only 40% DR. Claude Sonnet 4 is the *only* system achieving non-zero DR across all six domains, with the highest category alignment (CA = 0.489) of any evaluated system, though FPR escalates from 0% on SWE-BENCH to 67% on Toolathlon. No single system dominates all cells of the complete matrix, establishing that ensemble approaches combining structural and semantic detection are the most promising direction for comprehensive agent monitoring coverage.

1 Introduction

Large language model (LLM) agents are transitioning from research prototypes to production systems that execute consequential business workflows: customer service interactions, software engineering tasks,

financial analysis, and legal due diligence (ZenML, 2025). This transition brings an urgent need: when an agent fails—by executing the wrong action, producing incomplete analysis, or violating an operational policy—someone or something must detect the failure before it propagates.

The observability landscape has responded. Commercial platforms such as LangSmith, Langfuse, Arize Phoenix, and Braintrust now offer trace visualization, LLM-as-judge evaluation, and anomaly detection for agent workflows. Academic work has proposed moving beyond black-box benchmarking toward runtime analytics of agentic systems (Moshkovich et al., 2025). Yet a fundamental question remains unanswered: *how well do these monitoring tools actually work?*

Today there is no principled way to answer this question. Existing benchmarks evaluate agent *capability*—whether GPT-4 can resolve a GitHub issue (Jimenez et al., 2024), complete a customer service task (Yao et al., 2025), or navigate a web environment (Zhou et al., 2024). These benchmarks produce agent trajectories and ground-truth outcomes, but they do not evaluate whether a monitoring tool can correctly detect failures *within* those trajectories. The distinction matters: a monitoring tool must not only determine that something went wrong but must also identify *what* went wrong, *where* in the trajectory it occurred, and *how severe* the failure is.

The gap is consequential. Production deployments of multi-agent systems report failure rates between 41% and 87% (ZenML, 2025; Cemri et al., 2025). AI safety incidents increased 56.4% from 2023 to 2024 (Responsible AI Labs, 2025). Enterprise adoption demands reliability, but reliability requires monitoring, and monitoring tools have no benchmark against which to be measured.

Why This Benchmark Is Necessary

The cost of undetected agent failures is not hypothetical. In July 2025, Replit’s AI coding agent—operating during a declared code freeze—executed a `DROP TABLE` command that destroyed months of production data for over 1,200 users, then fabricated approximately 4,000 fake records in an apparent attempt to conceal the destruction.¹ A monitoring tool watching for destructive database operations during a declared freeze period would have flagged this trajectory before the irreversible action executed.

In the legal domain, attorneys submitted court filings containing entirely fabricated case citations generated by ChatGPT (*Mata v. Avianca*, S.D.N.Y., 2023), resulting in sanctions. Deloitte’s M&A advisory practice documented systematic quality failures in AI-generated due diligence reports that covered financial metrics thoroughly while omitting or providing cursory analysis of regulatory risk, antitrust exposure, and data privacy compliance (Deloitte Switzerland, 2025). These are precisely the kind of *quality degradation* failures—technically complete outputs that miss critical factors—that existing monitoring tools cannot detect because they lack domain-specific quality rubrics.

Customer-facing deployments have produced equally consequential failures. Air Canada’s chatbot fabricated a bereavement fare policy, advising a passenger he could retroactively apply for a discount that did not exist; the airline was held legally liable for the chatbot’s statements (Canadian Small Claims Court, 2024). New York City’s MyCity chatbot advised business owners to violate cash-acceptance and anti-discrimination laws. In each case, the agent trajectory contained clear signals—hallucinated policies, responses contradicting system prompts—that a properly instrumented monitoring tool could have detected.

These incidents span our full failure taxonomy: safety violations (Replit), hallucinated facts (*Mata v. Avianca*, Air Canada), quality degradation (Deloitte M&A), and policy violations (NYC MyCity). Each was detectable from the agent’s trajectory alone—the tool calls, the system prompts, and the agent’s responses contained sufficient signal. What was missing was a monitoring layer that could analyze these trajectories automatically. ATFD exists to measure whether such monitoring tools actually work.

¹Fortune, July 2025. “AI coding tool Replit wiped database, called it a catastrophic failure.”

We address this gap with ATFD (Agent Trajectory Failure Detection), a benchmark that evaluates monitoring tools rather than agents. Our contributions are:

1. **Failure taxonomy.** A 7-category, 23-subcategory taxonomy of agent trajectory failures (§3), grounded in prior failure analysis work (Cemri et al., 2025; Winston et al., 2025; Microsoft, 2025) and extended with *quality degradation* as a novel first-class category covering shallow outputs, suboptimal approaches, and incomplete analysis.
2. **Task definition.** A formal specification of the failure detection task (§4): given a complete trajectory, produce structured findings with severity, failure category, and step-level attribution, along with a mandatory cost report.
3. **Multi-domain dataset.** A corpus of 2,462 trajectories (§6) spanning customer service (τ -bench retail, airline, and telecom), coding (SWE-BENCH, 100 OpenHands trajectories), hand-crafted synthetic scenarios covering safety, quality, and infrastructure failures, ATBench (200-sample subset of 1,000 agent safety trajectories with 14 failure modes and 8 risk sources), and Toolathon (100 long-horizon multi-tool trajectories from HKUST-NLP, ICLR 2026).
4. **Ground truth protocol.** A multi-source consensus mechanism (§7) combining programmatic verdicts (test suites, state checks) with three independent LLM judges, validated by inter-rater agreement metrics.
5. **Complete comparative evaluation.** A full 7-system $\times$ 6-domain evaluation (§8–§9)—a naive heuristic baseline, the Galea zero-config structural heuristic, two open-weight LLM judges (Llama 3.3 70B on synthetic only due to free-tier rate limits; Llama 4 Scout across synthetic, retail, airline, and SWE-BENCH), Claude Sonnet 4 as a proprietary LLM judge across all six domains, LangSmith with custom evaluator rules, and Braintrust with custom scorers—measured on detection rate, false positive rate, F1, category alignment, cost, and configuration effort. The complete matrix reveals that no single system dominates all cells. Claude Sonnet 4 achieves the highest F1 on real retail trajectories (0.800) and the highest category alignment (CA = 0.489), but reaches only 40% DR on airline—the hardest domain for LLM judges—and 67% FPR on Toolathon. ATBench shows the starkest heuristic-vs-LLM gap: heuristic systems achieve 0% DR while Claude Sonnet 4 achieves 100% DR. Toolathon is the hardest domain for all systems except Claude Sonnet 4: every other system achieves 0% DR on long-horizon multi-tool trajectories, confirming that multi-step plan failures require semantic understanding of tool sequencing that structural and rule-based approaches cannot provide. LangSmith and Braintrust both achieve 100% DR on SWE-BENCH but 100% FPR, confirming that event-count thresholds are unusable without calibration. Braintrust additionally reveals a domain calibration problem: its threshold achieves 91.7% DR on retail but 71.1% FPR.
6. **Open benchmark harness.** A reproducible evaluation framework with cost-aware metrics, confidence intervals, and statistical tests, released as open-source software.

2 Related Work

ATFD draws on and extends four streams of prior work: agent benchmarks, LLM evaluation methods, agent failure analysis, and runtime monitoring.

2.1 Agent Benchmarks

A rich set of benchmarks evaluates LLM agent *capability*. τ -bench (Yao et al., 2025) simulates customer-service conversations in retail and airline domains, introducing the pass^k reliability metric that reveals

how single-trial results mask unreliability. τ^2 -bench (Research, 2025) extends this with a banking domain and quality fixes. SWE-BENCH (Jimenez et al., 2024) evaluates agents on 2,294 real GitHub issues with automated test-suite ground truth. AgentBench (Liu et al., 2024) spans 8 diverse environments and identifies poor long-term reasoning as the dominant failure mode. AgentBoard (Ma et al., 2024) introduces fine-grained progress-rate metrics that motivate ATFD’s three-tier outcome (pass/degraded/fail). WebArena (Zhou et al., 2024) provides realistic web environments where best agents achieve only 10–14% task success. GAIA (Mialon et al., 2024) demonstrates a dramatic gap between human (92%) and AI (15%) performance on real-world multi-step tasks. ToolLLM (Qin et al., 2024) evaluates tool use across 16,000+ APIs and shows that argument errors and tool selection are dominant failure modes.

All of these benchmarks share a common limitation from ATFD’s perspective: they evaluate the *agent* but not the *monitoring tool* that observes the agent. ATFD repurposes their trajectories as input to a meta-evaluation: given a trajectory that τ -bench or SWE-BENCH has labeled as pass or fail, can a monitoring system correctly identify the failure mode?

2.2 LLM Evaluation Methods

The LLM-as-judge paradigm, established by MT-Bench and Chatbot Arena (Zheng et al., 2023), provides the methodological foundation for using LLM verdicts as evaluation signals. Zheng et al. (2023) demonstrate that GPT-4 achieves >80% agreement with human experts but exhibits position bias, verbosity bias, and self-enhancement bias. G-Eval (Liu et al., 2023) shows that structured chain-of-thought rubrics substantially improve LLM-judge alignment with human judgments—a finding that directly informs ATFD’s domain-specific quality rubrics. Length-Controlled AlpacaEval (Dubois et al., 2024) addresses length bias, a concern for ATFD since verbose-but-incorrect trajectories may receive inflated scores from naive judges. Gu et al. (2024) catalog at least 8 distinct bias types and recommend multi-judge consensus as mitigation—exactly the approach ATFD adopts for ground truth construction.

2.3 Agent Failure Analysis

Several concurrent efforts develop failure taxonomies for LLM agents. MAST (Cemri et al., 2025) introduces a 14-failure-mode taxonomy across 3 categories validated on 1,600+ traces, finding that specification ambiguity and unstructured coordination account for 79% of multi-agent failures. ATFD’s taxonomy extends MAST by covering single-agent scenarios, distinguishing communication from state failures, and adding quality degradation. Anonymous (2025d) analyze 900 traces and identify failure to recover from tool errors as the dominant pattern, with suboptimal strategies as a distinct mode—mapped to ATFD’s `quality.suboptimal_approach`. Agent-SafetyBench (Zhang et al., 2024) evaluates 8 safety risk categories with 2,000 test cases. Winston et al. (2025) propose a tool-failure taxonomy with categories directly mapped to ATFD’s action subcategories. Microsoft’s industry whitepaper (Microsoft, 2025) and AGENTS SAFE (Anonymous, 2025a) provide complementary perspectives from industry governance. Vadlamudi (2026) independently proposes four functional dimensions that validate ATFD’s category structure. ToolFuzz (ETH SRI, 2025) demonstrates systematic tool failure injection, informing ATFD’s synthetic dataset generation.

The foundational framing of AI safety problems by Amodei et al. (2016)—avoiding side effects, reward hacking, and distributional shift—provides the intellectual lineage for treating safety failures as a distinct category. OWASP’s Agentic AI Top 10 (OWASP Foundation, 2025) and AgentLeak (Anonymous, 2026a) provide practitioner-oriented risk catalogs. Real-world incident data shows AI safety incidents increased 56.4% year-over-year, with hallucinations accounting for 38% (Responsible AI Labs, 2025).

2.4 Runtime Monitoring and Agent Observability

The closest prior work to ATFD’s contribution is [Moshkovich et al. \(2025\)](#), who argue for moving beyond black-box benchmarking toward observability-driven analytics of agentic systems. Their paper proposes a framework for how such tools should work; ATFD provides a concrete benchmark for evaluating whether they do.

Agent trajectory analysis is a specific instance of *process mining*, the discipline of extracting insights from event logs ([van der Aalst, 2016](#)). ATFD’s failure detection task is analogous to conformance checking: comparing an observed trajectory against expected behavior. [Berti et al. \(2025\)](#) demonstrate that LLMs can perform conformance checking with appropriate prompting, validating ATFD’s LLM-judge baseline design.

Broader surveys of agent evaluation ([Anonymous, 2025f,e](#)) confirm the gap ATFD addresses: no existing work evaluates the performance of monitoring tools on agent trajectories. The multi-dimensional evaluation framework CLEAR ([Anonymous, 2025b](#)) motivates ATFD’s inclusion of cost as a first-class metric. Practitioner analyses from Arize ([Arize AI, 2024](#)), Inkeep ([Inkeep, 2025](#)), and ZenML ([ZenML, 2025](#)) document the production failure landscape that monitoring tools must address. The enterprise reliability gap ([Simmering, 2025](#)) underscores the need for reliability-oriented benchmarks.

Methodological foundations for ATFD’s evaluation include Wilson score intervals ([Wilson, 1927](#)) for proportional metrics, bootstrap methods ([Efron and Tibshirani, 1993](#)) for aggregate confidence intervals, Fleiss’ κ ([Fleiss, 1971](#)) for inter-rater agreement, and McNemar’s test ([McNemar, 1947](#)) for pairwise system comparison. Principled selection of agreement metrics follows the recommendations of [Anonymous \(2026b\)](#), while [Anonymous \(2025c\)](#) caution that high agreement does not guarantee correct labels.

3 Agent Trajectory Failure Taxonomy

We propose a failure taxonomy comprising 7 top-level categories and 23 subcategories (Table 1). The taxonomy is designed to be *exhaustive* (every observed trajectory failure maps to at least one subcategory), *non-overlapping at the category level* (categories represent distinct failure dimensions), and *actionable* (each subcategory suggests a specific remediation strategy).

Relationship to prior taxonomies. The action, process, and safety categories overlap substantially with MAST ([Cemri et al., 2025](#)), Microsoft’s taxonomy ([Microsoft, 2025](#)), and the tool-failure taxonomy of [Winston et al. \(2025\)](#). ATFD’s primary extension is the `quality` category with 5 subcategories. Prior work treats quality degradation as “technically correct”—and therefore not a failure. We argue that in production deployments, a shallow analysis that covers 2 of 8 relevant risk factors *is* a failure, even if the agent did not crash, call the wrong tool, or violate a policy. The `quality` category captures this class of failures that matter to practitioners ([Arize AI, 2024](#); [Inkeep, 2025](#)) but are missed by binary pass/fail evaluation.

Taxonomy validation. We validated the taxonomy against three data sources: (1) failure annotations from τ -bench and SWE-BENCH trajectories, (2) the MAST failure mode catalog ([Cemri et al., 2025](#)), and (3) practitioner incident reports from Arize ([Arize AI, 2024](#)) and Responsible AI Labs ([Responsible AI Labs, 2025](#)). Every failure in these sources maps to at least one subcategory; no source contained failures requiring a new top-level category.

4 Task Definition

The ATFD task is defined as follows.

Input. A complete agent trajectory $T = (e_1, e_2, \dots, e_n)$ consisting of n events. Each event e_i has a type (user message, assistant message, tool call, tool result, or system event), a timestamp, content, and optional metadata. The trajectory also includes a natural language task description.

Output. A judge output $J(T)$ consisting of:

1. A boolean `has_failure` flag indicating whether the trajectory contains any failure.
2. A list of *findings*, each specifying:
 - **Severity:** `error`, `warning`, or `info`.
 - **Category:** a dotted string from the taxonomy (e.g. `action.wrong_args`, `quality.shallow_output`).
 - **Description:** a natural language explanation.
 - **Attribution** (optional): the step index or event range in the trajectory where the failure occurred.
3. A *cost report* specifying dollar cost, latency in seconds, total tokens consumed, number of API calls, and infrastructure tier (none, API key, hosted service, or GPU required).

Three-tier outcome model. Following AgentBoard’s insight that binary pass/fail metrics mask meaningful performance variations (Ma et al., 2024), ATFD defines three ground truth outcomes:

- **Pass:** the trajectory completes the task correctly with acceptable quality.
- **Degraded:** the trajectory produces a technically correct result but with substandard quality (maps to the `quality` category).
- **Fail:** the trajectory contains a hard failure in one or more of the remaining 6 categories.

This three-tier model enables separate measurement of hard-failure detection (Detection Rate, DR) and quality-degradation detection (Quality Detection Rate, QDR), revealing monitoring tools’ sensitivity to subtle failures.

5 Metrics

ATFD evaluates monitoring systems on seven metrics, all reported with 95% confidence intervals.

Detection Rate (DR). Recall on `fail` trajectories: the proportion of ground-truth failures for which the system raised at least one `error` or `warning` finding.

$$\text{DR} = \frac{|\{T : \text{gt}(T) = \text{fail} \wedge \text{detected}(T)\}|}{|\{T : \text{gt}(T) = \text{fail}\}|} \quad (1)$$

Quality Detection Rate (QDR). Recall on `degraded` trajectories: the proportion for which the system raised at least one `quality.*` finding.

$$\text{QDR} = \frac{|\{T : \text{gt}(T) = \text{degraded} \wedge \text{quality_detected}(T)\}|}{|\{T : \text{gt}(T) = \text{degraded}\}|} \quad (2)$$

False Positive Rate (FPR). The proportion of `pass` trajectories for which the system erroneously raised an `error-severity` finding.

$$\text{FPR} = \frac{|\{T : \text{gt}(T) = \text{pass} \wedge \text{error_raised}(T)\}|}{|\{T : \text{gt}(T) = \text{pass}\}|} \quad (3)$$

F1 Score. The harmonic mean of precision and recall for binary failure detection (`fail` vs. `pass/degraded`).

Category Alignment (CA). Macro-averaged F1 across all failure categories, treating each trajectory as a multi-label classification instance. This measures whether the system not only detects *that* a failure occurred but correctly identifies *what type* of failure it is.

Configuration Effort. A qualitative ordinal score (1–5) measuring the setup burden: API key only, prompt engineering required, domain-specific configuration, custom integration, or full pipeline development.

Cost. Mean dollar cost per trajectory, total cost for the full dataset, median latency (p50), and 95th-percentile latency (p95).

Confidence intervals. All proportional metrics (DR, QDR, FPR) are reported with Wilson score intervals (Wilson, 1927), which provide accurate coverage even for proportions near 0 or 1 and for small samples. Aggregate metrics (macro-F1, cost means) use bootstrap confidence intervals (Efron and Tibshirani, 1993) with 10,000 resamples. Pairwise system comparisons use McNemar’s test with continuity correction (McNemar, 1947). Ground truth inter-rater agreement is measured via Fleiss’ κ (Fleiss, 1971).

6 Datasets

The ATFD benchmark draws trajectories from five sources (Table 2).

τ -bench trajectories. We draw from τ -bench (Yao et al., 2025), using the retail (114 tasks $\times$ 4 trials = 456 trajectories), airline (50 tasks $\times$ 4 trials = 200 trajectories), and telecom (114 tasks $\times$ 4 trials = 456 trajectories) domains. Ground truth is derived from τ -bench’s state-comparison mechanism: each task specifies expected database mutations, and `pass/fail` is determined by comparing actual vs. expected state. Observed failure rates are 22% (retail), 28% (airline), and 41% (telecom). Extension to the three-tier outcome via domain-specific quality rubrics (§7) is planned as future work.

SWE-BENCH trajectories. We include 100 trajectories from the `nebius/SWE-rebench-openhands-trajectories` dataset on HuggingFace (Jimenez et al., 2024), generated by OpenHands v0.54.0. The dataset is balanced: 50 resolved (`pass`) and 50 unresolved (`fail`) trajectories. Ground truth is provided by the SWE-BENCH automated test suite: a trajectory is labeled `pass` if all required tests pass after the agent’s patch, and `fail` otherwise. Unlike τ -bench, SWE-BENCH trajectories involve file editing, code generation, bash command execution, and test execution, exercising structurally different failure modes: unresolved trajectories typically contain explicit tool errors (bash failures, test failures) as the agent attempts and fails to build or test its solution. This structural obviousness makes SWE-BENCH the *easy* domain for monitoring tools (see §11.5).

Synthetic trajectories. We hand-craft 50 trajectories (100% failure rate by design) to ensure coverage of underrepresented failure categories—particularly safety (`permission_escalation`, `data_leakage`) and quality (`poor_tone`, `low_confidence_output`)—that rarely occur naturally in benchmark trajectories. Synthetic trajectories are constructed using templates with injected failures, following the methodology of ToolFuzz (ETH SRI, 2025) and Agent-SafetyBench (Zhang et al., 2024).

ATBench trajectories. We include a 200-sample subset from the AI45Research/ATBench dataset on HuggingFace (AI45Research, 2025) (full dataset: 1,000 trajectories). ATBench covers agent safety scenarios with 14 failure modes and 8 risk sources; the full dataset is nearly balanced (503 safe, 497 unsafe), and we sample proportionally (approximately 101 safe, 99 unsafe) to preserve this balance. Failure modes span tool description injection, harmful action execution, data leakage, and related agent safety violations. Ground truth is provided by the ATBench evaluation protocol. Unlike structural failure domains, ATBench failures are *semantically* defined: no structural signal (tool errors, loop patterns, step counts) differentiates safe from unsafe trajectories. This makes ATBench the hardest domain for heuristic systems and confirms a fifth tier of the domain difficulty spectrum in which structural analysis is categorically insufficient (see §11.6).

Toolathlon trajectories. We include 100 trajectories from the `hkust-nlp/Toolathlon-Trajectories` dataset on HuggingFace (HKUST-NLP, 2026), generated by Claude 4 Sonnet on long-horizon multi-tool tasks spanning 600+ available tools. Toolathlon (HKUST-NLP, ICLR 2026) is distinguished by the scale and complexity of its trajectories: individual runs may contain up to 16,000 messages across hundreds of tool calls. We cap each trajectory at 100 events for evaluation to bound cost and latency; the full trajectory is retained for statistical analysis. Ground truth is derived from the `task_status.evaluation` field: trajectories labeled `true` (pass) or `false` (fail) are included; `null` (inconclusive) is excluded. Unlike τ -bench, where failures involve argument or state errors, and unlike ATBench, where failures are semantically clear harmful instructions, Toolathlon failures involve *multi-step plan failures*: the agent selects a coherent but incorrect sequence of tools, reaches an intermediate dead end, or misunderstands dependencies across a long context. This makes Toolathlon the hardest domain for all evaluated systems except Claude Sonnet 4 and establishes a sixth tier of the domain difficulty spectrum (see §11.7).

7 Ground Truth Validation

Establishing reliable ground truth for agent trajectory failures is challenging: unlike classification tasks with objective labels, failure detection involves subjective judgments about severity and category boundaries. We adopt a multi-source consensus protocol.

Source 1: Programmatic verdicts. For τ -bench trajectories, we use the benchmark’s built-in state comparison. For SWE-BENCH, we use the automated test suite verdicts embedded in the `nebius/SWE-rebench-openhands-t` dataset: each trajectory is pre-labeled resolved (pass) or unresolved (fail) based on whether the agent’s patch causes all required tests to pass. These provide high-confidence binary (pass/fail) labels but do not identify failure categories or quality degradation.

Source 2: LLM judge consensus (planned). In the planned full protocol, three independent LLM judges—GPT-4.1, Claude Sonnet 4, and Llama 4 Scout—will each analyze every trajectory using a structured two-stage prompt. Stage 1 asks the judge to identify all potential issues; Stage 2 asks the judge to classify each issue using the ATFD taxonomy with severity assignments. In this release, we use three LLM judges (Llama 3.3 70B, Llama 4 Scout, and Claude Sonnet 4) as evaluated systems; full multi-judge consensus across all trajectories is pending.

Source 3: Domain-specific quality rubrics. For degraded labeling, we apply domain-specific quality rubrics inspired by G-Eval’s form-filling paradigm (Liu et al., 2023). Each domain defines 3–5 quality dimensions scored on a 0–2 scale (adequate, marginal, inadequate). A trajectory is labeled degraded when the programmatic check passes but the average quality score falls below a domain-specific threshold.

Consensus mechanism. The final ground-truth label is determined by majority vote across sources. A label is assigned gold consensus when all sources agree, majority when ≥ 3 of 4 sources agree, and disputed otherwise. Disputed trajectories are retained with their majority label but flagged for analysis.

Agreement metrics. Inter-rater agreement among the LLM judges is measured via Fleiss’ κ (Fleiss, 1971) on the three-tier outcome. The full multi-judge consensus pipeline (GPT-4.1, Claude Sonnet 4, and Llama 4 Scout) has not yet been executed; we plan to report $\kappa \geq 0.7$ as the acceptance threshold following Anonymous (2026b). In the current evaluation, ground truth is derived solely from programmatic verdicts (τ -bench state comparison for real trajectories, injected labels for synthetic), with LLM judge consensus validation planned as future work. As cautioned by Anonymous (2025c), we will validate consensus labels against programmatic verdicts where available to confirm that agreement reflects accuracy, not systematic bias.

8 Evaluated Systems

We evaluate 7 systems in this release and describe 1 additional planned system (Table 4).

Naive Heuristic. A rule-based baseline that checks for error strings in tool results, timeout signals, step-count thresholds, and repeated identical tool calls. This system cannot detect quality degradation, communication failures, or most safety failures. It serves as a zero-cost lower bound.

Galea Heuristic. The Galea heuristic is a zero-configuration structural analysis system that requires no LLM calls and no domain-specific setup. It applies a richer set of pattern-matching rules than the naive heuristic: tool-error fingerprinting across multiple error taxonomies, loop and cycle detection in tool-call sequences, state-signal analysis (comparing expected vs. observed state transitions), and structural anomaly detection in turn-taking patterns. Galea’s heuristic makes no calls to any language model and incurs \$0.00 cost per trajectory. Latency is sub-100ms (p50: 39–54ms across datasets).

LLM Judges (evaluated). Two open-weight LLM judges—Llama 3.3 70B and Llama 4 Scout—accessed via Groq’s free API tier. Each receives the full trajectory and a system prompt describing the ATFD taxonomy, instructing structured JSON output with findings. Both use identical prompts to isolate model capability differences. The choice of free-tier models is deliberate: cost is a first-class metric, and demonstrating what is achievable at zero marginal cost establishes a practical lower bound for LLM-based monitoring. Llama 4 Scout was evaluated on all six domains (n=50 synthetic, n=50 retail, n=20 airline, n=20 SWE-BENCH, n=200 ATBench, n=100 Toolathlon). Llama 3.3 70B was evaluated only on synthetic (n=50): Groq free-tier rate limits made all other domains infeasible within the evaluation window.

Claude Sonnet 4 Judge (evaluated). Claude Sonnet 4 was accessed via the Anthropic API and evaluated across all six domains: synthetic (n=50), retail (n=20), airline (n=20), SWE-BENCH (n=20), ATBench (n=200), and Toolathlon (n=100). It receives the same full trajectory and taxonomy-structured system prompt as the Llama judges. Claude Sonnet 4 is the only system achieving non-zero DR across all six domains:

100% DR on synthetic, retail, SWE-BENCH, ATBench, and Toolathlon, with 40% DR on airline. It delivers the highest category alignment (CA = 0.489 on synthetic) of any evaluated system and is the only LLM judge achieving both 100% DR and 0% FPR on SWE-BENCH. Its FPR escalates on harder domains: 6% on retail, 20% on airline, and 67% on Toolathlon—the last reflecting calibration challenges in extremely long-context multi-tool trajectories.

LangSmith (evaluated). LangSmith is a commercial observability platform for LLM applications that supports custom evaluator rules written as Python functions. We configured 4 evaluation rules (~80 lines of code) targeting the most tractable failure categories: a tool-loop detector (triggering on repeated identical tool calls), and an event-count threshold that flags runs exceeding a configurable step budget. Setup required approximately 15 minutes of configuration including account creation, project setup, and rule authoring. LangSmith’s evaluation model is fundamentally *rule-driven*: each failure category requires explicit authorship, in contrast to LLM judges (which generalize from a prompt) or the Galea heuristic (which ships zero-config). This places LangSmith in a distinct configuration tier (Effort = 4) reflecting substantial per-deployment setup burden.

Braintrust (evaluated). Braintrust is a commercial evaluation platform for LLM applications that supports custom scorer functions written in JavaScript or Python. We configured the identical 4 failure-detection rules used for LangSmith (approximately 60 lines of code in Braintrust’s scorer format), targeting the same failure categories: tool-loop detection and an event-count threshold. Setup required approximately 15 minutes including account creation, project setup, and scorer authorship. Crucially, this configuration replicates LangSmith’s rules exactly, enabling a direct comparison of rule-driven platforms when holding the rule set constant.

Planned systems (not yet evaluated). One additional system is planned for future evaluation: GPT-4.1 judge (requiring paid OpenAI API access). This system is listed in Table 4 but excluded from all results tables. LangSmith and Braintrust are both fully evaluated across all six domains in this release (see Table 5).

9 Results

9.1 Main Results

Table 5 presents the complete evaluation results for all seven evaluated systems across all six benchmark domains. Quality Detection Rate (QDR) is omitted because the current dataset uses binary pass/fail labels; the *degraded* tier requires quality rubrics not yet applied. Note that sample sizes vary by cell: LLM judges were evaluated on subsets due to API throughput constraints, while heuristic systems ran on full dataset splits. Llama 3.3 70B was evaluated on synthetic only; Groq free-tier rate limits prevented all other domain runs (see §8).

Summary matrix (DR % / FPR %)

Table 6 condenses the complete results into a single compact view for cross-system, cross-domain comparison.

9.2 Per-Domain Breakdown

The complete matrix (Tables 5 and 6) enables seven cross-domain observations that are only visible once all cells are populated.

- **Domain difficulty spectrum is clear and now six-tiered.** SWE-BENCH is the easiest domain (all systems achieve high DR on structural coding errors), synthetic is moderate (LLM judges excel, pure heuristics plateau at 78%), the τ -bench domains are hard—with airline harder than retail for LLM judges (Llama 4 Scout: 20% vs. 25%; Claude Sonnet 4: 40% vs. 100%)—ATBench reveals a middle ground for heuristics (Galea: 52.5% DR, between its 0% on Toolathlon and 100% on τ -bench), and Toolathlon is the hardest domain for all systems except Claude Sonnet 4. The spectrum: SWE-BENCH $\rightarrow$ Synthetic $\rightarrow$ Retail $\rightarrow$ Airline $\rightarrow$ ATBench $\rightarrow$ Toolathlon, each posing qualitatively distinct challenges.
- **Rule-based platforms collapse on SWE-BENCH.** LangSmith and Braintrust both achieve 100% DR on SWE-BENCH—but with 100% FPR. Their event-count threshold triggers on every passing coding trajectory, making the configuration operationally unusable on this domain without domain-specific calibration.
- **Airline is the hardest domain for LLM judges.** Claude Sonnet 4 reaches only 40% DR on airline (vs. 100% on retail); Llama 4 Scout reaches only 20% (vs. 25% on retail). Airline failures involve subtle state errors and policy interpretations that confuse even frontier LLMs. In contrast, the Galea heuristic achieves 100% DR on airline with 0% FPR.
- **Toolathlon is hardest for all systems except Claude Sonnet 4.** Every system except Claude Sonnet 4 achieves exactly 0% DR on Toolathlon: the naive heuristic, Galea, LangSmith, and Braintrust all emit no findings on long-horizon multi-tool trajectories; Llama 4 Scout reaches only 14.3% DR with 0% FPR. Only Claude Sonnet 4 achieves 100% DR on Toolathlon, but at a high FPR (66.7%)—the highest FPR of any Claude cell across all domains. Long-horizon failures require semantic understanding of multi-step plans; structural pattern matching and simple rules cannot detect whether a 16,000-message trajectory with hundreds of tool calls reached the correct goal.
- **ATBench reveals a middle ground for Galea.** Unlike prior domains where Galea achieves either 100% DR (τ -bench) or 0% DR (Toolathlon), Galea achieves 52.5% DR on ATBench with 0% FPR (error-severity definition). This confirms that some safety failure patterns—tool-description anomalies and structural irregularities that co-occur with unsafe actions—are partially detectable by structural analysis, while the remaining 47.5% require semantic understanding.
- **Claude Sonnet 4 is the only system with 100% DR across all six domains**, though FPR escalates on harder domains (0% on SWE-BENCH, 6% on retail, 20% on airline, 67% on Toolathlon). Its FPR on retail (5.6%) is substantially lower than Llama 4 Scout (13.2%) and the naive heuristic (17.8%). Its category alignment (CA = 0.489 on synthetic) remains the highest of any evaluated system.
- **Galea is the only system with 0% FPR across all domains** (using the error-severity definition). Galea achieves 100% DR on both retail and airline with 0% FPR, 90% on SWE-BENCH, 52.5% on ATBench, and 78% on synthetic—in every case with 0% FPR. Its 0% DR on Toolathlon is the one cell where structural analysis is categorically insufficient.
- **No system dominates all cells.** The matrix reveals complementary strengths: Galea dominates on real τ -bench domains (100% DR, 0% FPR on retail and airline); Claude Sonnet 4 is the only system achieving 100% DR on all six domains; LLM judges and rule-driven platforms uniformly achieve 100% DR on SWE-BENCH but vary wildly on FPR. This complementarity is the central empirical finding: an ensemble combining structural and semantic detection would outperform any single system on the complete six-domain matrix.

[Figure pending: per-category detection rates will be generated once per-category annotations are available for τ -bench trajectories and additional LLM judges are evaluated.]

Figure 1: Per-category detection rates (planned). Currently, per-category breakdown is available only for the 50-trajectory synthetic dataset.

9.3 Per-Category Analysis

Per-category detection rates are currently available only for the synthetic dataset, where injected failure types are known by construction (Table 3). On synthetic data, the LLM judges achieve 100% detection across all injected categories (process, communication, safety), while the naive heuristic detects only 14% overall—primarily `tool_loop` patterns.

Category alignment (CA) scores tell the more nuanced story: even when LLM judges correctly detect that a failure exists, open-weight judges achieve only CA = 0.359 (Llama 3.3 70B) and CA = 0.318 (Llama 4 Scout) on synthetic data, while Claude Sonnet 4 achieves the highest CA of any system at CA = 0.489 on synthetic data—indicating that identifying *what* failed is genuinely harder than detecting *that* a failure occurred, even for a frontier model. On real τ -bench data, CA drops to 0.133 (Claude Sonnet 4 on retail), 0.029 (Llama 4 Scout on retail), and 0.000 (naive heuristic on all real datasets).

Key observations:

- **Process** failures (`tool_loop`, `context_overflow`) are the most reliably detected category in synthetic data, as they produce structural patterns.
- **Safety** failures (`permission_escalation`, `data_leakage`) are detected by LLM judges in synthetic scenarios but require explicit policy context.
- **Communication** failures (`hallucination`) are detected at 100% in synthetic data where the hallucination is egregious, but real-world subtle hallucinations remain untested.
- **Quality** failures are not represented in the current synthetic set and cannot be assessed on τ -bench without quality rubrics—this remains the frontier challenge.

9.4 Cost Analysis

Table 7 summarizes the cost profile of each evaluated system. A key finding is that both LLM judges were accessed via Groq’s free API tier, making their marginal dollar cost \$0.00. Cost manifests instead as *rate limits*: evaluating 456 retail trajectories with Llama 4 Scout was infeasible under free-tier rate limits, forcing us to evaluate only a 50-trajectory subset. This demonstrates that for practical deployment, cost should be measured not just in dollars but in *throughput constraints*.

10 Ablation Study

Full ablation studies are planned for future work, contingent on access to the additional systems and completion of the full evaluation matrix. We outline the planned ablations here to guide future investigation.

10.1 Naive Heuristic Ablation (Planned)

The naive heuristic system comprises four independent rule groups: error string matching, step-count thresholds, loop detection, and timeout detection. We plan to ablate each group individually to quantify its contribution to overall DR. Preliminary observations suggest that the high DR on airline (40.9%) is driven primarily by error string matching, while loop detection contributes most to synthetic trajectory detection.

10.2 LLM Judge Prompt Ablation (Planned)

We plan to ablate the LLM judge prompt by removing components (taxonomy reference, structured output format, domain context) to measure their individual contributions to DR and CA. The low CA scores observed across all systems (0.000–0.359) suggest that the taxonomy reference in the prompt may not be sufficient for accurate category assignment, and alternative prompting strategies (*e.g.* chain-of-thought, few-shot examples) warrant investigation.

11 Analysis

11.1 Complementary Detection Profiles

The complete evaluation matrix reveals that heuristic and LLM-based systems exhibit *inverse* detection profiles across domains: heuristic systems excel at structural failures in real trajectories, while LLM judges excel at semantic failures in synthetic trajectories. Table 8 summarizes this complementarity with the full six-domain view.

The Galea heuristic achieves 100% DR on both retail and airline τ -bench splits with 0% FPR because these failures predominantly involve tool-error fingerprints, state-signal anomalies, and loop patterns—all structurally detectable without language understanding. The 78% DR on synthetic data reveals the ceiling of pure structural analysis: the 22% missed failures are communication failures (*hallucination*) and safety failures (*data_leakage*, *permission_escalation*) that produce no structural anomaly and require semantic grounding to identify.

Open-weight LLM judges (Llama 3.3 70B, Llama 4 Scout) detect 100% of synthetic failures—including hallucinations and safety violations—because the injected failures are egregiously obvious at the semantic level. Their performance on real τ -bench domains reveals the opposite gap: Llama 4 Scout achieves only 25% DR on retail and a worse 20% DR on airline, with a severe 67% FPR on airline. Natural agent failures involve subtle argument value errors, near-correct state transitions, and ambiguous intent that confuse semantic analysis trained on more salient failure patterns. Airline is harder than retail precisely because airline policies are more domain-specific and failures are less structurally obvious.

Claude Sonnet 4 stands apart by bridging structural and semantic profiles on *some* domains: it achieves 100% DR on both synthetic and real retail trajectories with $F1 = 0.800$ on retail—the best real-trajectory performance of any evaluated system. Its FPR on retail (5.6%) lies between the heuristics (0%) and open-weight judges (13.2%), and its CA of 0.489 on synthetic data is the highest observed across all systems. However, the complete matrix reveals that Claude Sonnet 4 does not generalize uniformly: its airline DR of 40% and FPR of 20% show that frontier proprietary models still struggle with domain-specific policy semantics in the airline domain, where Galea’s structural approach maintains superiority.

Ensemble implication. The complementarity across the complete six-domain matrix strengthens the case for an ensemble architecture: run the Galea heuristic first (zero latency, zero cost, zero FPR) to catch structural failures across all domains, then invoke Claude Sonnet 4 for synthetic, retail-style, ATBench, and Toolathlon trajectories where semantic understanding is needed. Neither system alone covers the full matrix; together

they would combine Galea’s 100% retail/airline coverage and 52.5% ATBench coverage with Claude’s 100% synthetic/SWE-BENCH/ATBench/Toolathlon coverage and superior category alignment—though Claude’s 67% FPR on Toolathlon remains an open challenge for ensemble precision. We leave formal ensemble evaluation as future work.

11.2 The Configuration Effort Gap

LangSmith’s results expose a third dimension of evaluation beyond detection rate and cost: *configuration effort*. With 4 manually authored evaluation rules (~80 lines of Python) and approximately 15 minutes of setup time, LangSmith achieves only 22.0% DR on the synthetic dataset—lower than the zero-config Galea heuristic (78.0% DR) and dramatically below the zero-config LLM judges (100% DR). Crucially, LangSmith’s 4 rules catch only what those rules explicitly target: tool-loop patterns (6 of 10 detected) and a coarse event-count threshold (5 of 50 flagged via step budget). Every other failure category—hallucination (0/10), permission_escalation (0/5), data_leakage (0/5), infinite_delegation (0/5), context_overflow (0/5), planning_failure (0/10)—is completely missed.

Platform-agnostic rule coverage. Braintrust was configured with the identical 4 rules (translated to Braintrust’s scorer format, ~60 LOC). It produces precisely the same 22.0% DR on the synthetic dataset. This equivalence is not coincidental: it proves that the detection gap is a property of the *rules*, not the *platform*. LangSmith and Braintrust are effectively interchangeable when given the same rule set; what differs is syntax, not capability. This finding validates the ATFD config effort metric: the bottleneck for rule-driven platforms is always the coverage of the authored rules, regardless of which platform hosts them.

This reveals a fundamental limitation of rule-driven platform monitoring: *coverage scales with authorship*. Catching a new failure category requires writing a new rule. For production deployments spanning dozens of failure modes, this becomes an ongoing engineering burden rather than a one-time setup cost.

Table 9 makes the configuration effort comparison explicit across the four evaluation paradigms represented in this study.

The practical implication is clear: monitoring tools that require per-category rule authorship create an ongoing maintenance cost as workflows evolve and new failure modes emerge. Zero-config approaches (heuristics, LLM judges) absorb this cost at development time rather than deployment time.

11.3 What Is Hard to Detect?

Our per-subcategory analysis reveals a consistent hierarchy of detection difficulty across systems:

Easy to detect (structural pattern matching). Tool errors, timeout signals, and tool-call loops in real τ -bench trajectories are detected at 100% by the Galea heuristic with 0% false positives. The naive heuristic detects only a fraction of these (25.4% on retail, 40.9% on airline), indicating that Galea’s richer pattern library—beyond simple error-string matching—captures a much broader range of structural failure signals.

Easy to detect (semantic understanding). Synthetic failures with obvious semantic patterns—egregious hallucination, explicit permission_escalation, clear tool_loop structures with labeled repetition—are detected at 100% by LLM judges. These are detectable because the failure signal is salient at the text level and the failure is designed to be identifiable.

Moderate difficulty. Action failures (wrong_tool, missing_action) require understanding task intent to determine correctness. LLM judges substantially outperform heuristics here. wrong_args is particularly challenging as it requires comparing argument values against task requirements.

Hard to detect. Quality failures are consistently the most difficult. `shallow_output` and `incomplete_analysis` require domain knowledge to assess whether the output is sufficiently comprehensive. `suboptimal_approach` requires understanding of what an efficient solution would look like. These findings corroborate practitioner reports that quality issues are underappreciated relative to hard crashes (Inkeep, 2025; Arize AI, 2024).

11.4 Domain Calibration Is Non-Trivial

The complete matrix exposes domain calibration as a pervasive problem for rule-driven platforms—not merely a retail anomaly. The event-count threshold (flagging runs with >20 events) was designed to catch runaway trajectories. Its behavior across all four domains reveals three distinct failure modes.

On the *synthetic* dataset, it contributes modestly to DR without FPR issues, because synthetic trajectories are short by construction. On *retail* τ -bench, the same threshold fires on 71.1% of passing trajectories: normal customer-service conversations routinely exceed 20 events. The result is 91.7% DR paired with 71.1% FPR—a configuration that would flood any production alerting system with noise. On SWE-BENCH, the calibration failure is even more severe: 100% DR with 100% FPR. Every passing coding trajectory triggers the threshold because software engineering tasks routinely involve dozens of tool calls (bash execution, file editing, test runs). The 100% FPR makes this configuration entirely unusable on SWE-BENCH without domain-specific recalibration. On *airline*, both LangSmith and Braintrust achieve 64% DR with 33% FPR—better calibrated than retail or SWE-BENCH, but still with a high false positive burden.

This four-domain view reveals a systematic pattern: *any* rule with a numeric threshold calibrated to one distribution will fail on a distribution with different event-count norms. Domain calibration is not a one-time tuning cost; it is an ongoing maintenance burden as new domains, agents, or workflows are added.

Zero-config systems avoid this failure mode by design. The Galea heuristic achieves 0% FPR across all four domains without threshold tuning because its structural patterns are data-distribution-invariant (a genuine tool error looks the same regardless of conversation length). LLM judges do not rely on numeric thresholds, though they introduce domain sensitivity via prompt-level assumptions—as visible in Claude Sonnet 4’s 20% FPR on airline despite 0% FPR on SWE-BENCH.

11.5 SWE-BENCH: The Easy Domain

SWE-BENCH trajectories present a structurally distinct failure profile that makes them qualitatively easier to monitor than τ -bench. Failed OpenHands trajectories from `nebius/SWE-rebench-openhands-trajectories` consistently exhibit explicit tool-level signals: bash command failures, compilation errors, and test suite failures appear directly in tool results. This structural obviousness allows even the naive heuristic to achieve 100% DR [92.9%, 100.0%] with 0% FPR on SWE-BENCH (n=100), compared to only 25.4% DR on the semantically subtle τ -bench retail domain.

The Galea heuristic achieves 90.0% DR [78.6%, 95.7%] on SWE-BENCH (n=100), slightly below the naive heuristic’s perfect score. This inversion—where Galea underperforms the naive baseline on a specific domain—reflects the structural nature of the task: tool-error string matching (the naive heuristic’s core mechanism) is exactly the right signal for SWE-BENCH, while Galea’s richer pattern library (loop detection, state-signal analysis) adds recall on more subtle domains but introduces no additional benefit here.

Claude Sonnet 4 on SWE-BENCH (n=20) achieves 100% DR [79.4%, 100.0%] with 0% FPR [0.0%, 16.8%]—the only LLM judge achieving both 100% DR and 0% FPR on this domain. Llama 4 Scout (n=20) and rule-driven platforms (n=100) also achieve 100% DR but with 100% FPR: every passing coding trajectory triggers the event-count threshold (rule platforms) or the LLM flags spurious issues in long tool-call sequences (Llama 4 Scout). This confirms that SWE-BENCH is “easy” for *detection* but *difficult for precision* for any system that relies on event-count heuristics or lacks calibrated confidence on passing trajectories.

Domain difficulty spectrum. The complete six-domain matrix reveals a clear difficulty ordering for monitoring tools:

- **SWE-BENCH (easy):** structural failures produce explicit error signals. All evaluated systems achieve 100% DR. The distinguishing challenge is precision: rule platforms and Llama 4 Scout achieve 100% FPR; only Claude Sonnet 4 and heuristics achieve both 100% DR and 0% FPR.
- **Synthetic (moderate):** injected failures span structural and semantic categories. LLM judges excel (100% DR); pure heuristics plateau at 78%; rule-driven platforms achieve only 22% DR.
- **Retail τ -bench (hard):** failures involve subtle argument errors and near-correct state transitions. Galea and Claude Sonnet 4 achieve high DR; open-weight judges and rule platforms struggle.
- **Airline τ -bench (hardest for LLM judges):** Even Claude Sonnet 4 reaches only 40% DR. Only Galea achieves both high DR (100%) and low FPR (0%) here.
- **ATBench (middle ground; partially structural, partially semantic):** Galea achieves 52.5% DR with 0% FPR—some safety patterns co-occur with structural anomalies that heuristics can detect. Claude Sonnet 4 achieves 100% DR and Llama 4 Scout achieves 70% DR, confirming that semantic understanding substantially outperforms structural analysis but is not uniquely necessary. Rule-driven platforms (LangSmith, Braintrust) achieve 0% DR, confirming that rules written for process/structural failures do not generalize to semantic safety.
- **Toolathlon (hardest for all systems except Claude Sonnet 4):** Long-horizon multi-tool failures involve no structural signal (no error strings, no loops, no step-count anomalies that are not normal for 600+-tool environments). Every system except Claude Sonnet 4 achieves 0% DR. Claude Sonnet 4 achieves 100% DR with 67% FPR: it can detect genuine long-horizon failures but over-flags passing trajectories, suggesting calibration challenges in extremely long contexts. Llama 4 Scout reaches 14.3% DR with 0% FPR, indicating that open-weight models can occasionally identify long-horizon failures but lack the consistent semantic depth of frontier models.

This six-tier spectrum has a critical practical implication: benchmarking monitoring tools on any single domain gives a highly misleading picture of production performance. The complete matrix—spanning structural-easy (SWE-BENCH) through long-horizon-hard (Toolathlon)—is necessary to characterize operational coverage across the full failure landscape. Notably, Claude Sonnet 4 is the only system achieving any non-zero DR across all six tiers, and no single system achieves high DR *and* low FPR uniformly.

Galea FPR: a metric asymmetry. The Galea heuristic reports 0% FPR across all evaluated domains, including SWE-BENCH. This is not accidental: Galea’s findings are emitted at `warning` severity by default, not `error` severity. The FPR metric (Eq. 3) counts only `error`-severity findings on passing trajectories. Because Galea never emits error-severity findings in its zero-config configuration, its FPR is structurally zero regardless of domain. The DR metric counts both `error` and `warning` findings, so Galea’s high DR reflects genuine structural pattern detection. Readers should interpret Galea’s 0% FPR as “no false error-severity alerts,” not as “zero false positives across all severity levels.” We report this distinction transparently to avoid inflating Galea’s apparent precision.

11.6 ATBench: Semantic Safety as the Inverse Domain

ATBench (AI45Research, 2025) provides an important counterpoint to the structural vs. semantic binary. Its 1,000 trajectories (503 safe, 497 unsafe) span 14 failure modes and 8 risk sources. We evaluate on a

proportionally sampled 200-trajectory subset (≈ 101 safe, ≈ 99 unsafe). Failure modes are defined at the semantic level: tool description injection, harmful action execution, data leakage, and related safety violations.

The complete ATBench results reveal a more nuanced picture than the original pilot suggested. Claude Sonnet 4 achieves 100% DR [98.1%, 100.0%] on the 200-trajectory evaluation, confirming that frontier models can reliably identify semantic safety failures. Llama 4 Scout achieves 70% DR on ATBench, indicating that open-weight models can detect a substantial majority of safety failures—a meaningful improvement over their performance on τ -bench domains. Most significantly, Galea achieves 52.5% DR [45.5%, 59.5%] on ATBench with 0% FPR: this is Galea’s *middle-ground* result, between its 100% on τ -bench and 0% on Toolathlon. Some ATBench safety patterns are structurally co-detectable (tool-description anomalies that deviate from normal schema patterns, unusual tool-call sequences preceding harmful outcomes), making partial structural detection possible.

Rule-driven systems (LangSmith, Braintrust) achieve 0% DR on ATBench, confirming that rules written for process and structural failures do not transfer to semantic safety.

ATBench in the difficulty spectrum. ATBench occupies a fifth tier in the domain difficulty spectrum—harder than airline for heuristic systems (0% for naive/rule-based; 52.5% for Galea) but easier than Toolathlon for LLM judges (70–100% vs. 0–14% on Toolathlon for open-weight models). Together with Toolathlon, ATBench completes the spectrum from structurally trivial (SWE-BENCH) through semantically hard (ATBench) to long-horizon impossible (Toolathlon).

Ensemble implication. The ATBench results strengthen the ensemble case: a Galea+Claude Sonnet 4 architecture would combine Galea’s 100% DR on τ -bench retail/airline (where Claude reaches only 40–100%) with Galea’s 52.5% structural coverage on ATBench and Claude’s 100% semantic coverage. Neither system alone covers the full six-domain matrix; together they would achieve comprehensive coverage across structural, semantic safety, and (partially) long-horizon failure classes—with Toolathlon remaining the hardest open problem.

11.7 Toolathlon: Long-Horizon Tool Use as the Hardest Domain

Toolathlon (HKUST-NLP, 2026) (HKUST-NLP, ICLR 2026) provides the most demanding evaluation in the ATFD benchmark. Its 100 trajectories—drawn from Claude 4 Sonnet runs on long-horizon multi-tool tasks with 600+ available tools—represent the frontier of agent complexity: individual runs may contain up to 16,000 messages, and task completion requires coordinating across dozens of sequential tool calls where intermediate outputs feed downstream steps.

The results establish Toolathlon as categorically different from all other evaluated domains. Every system except Claude Sonnet 4 achieves exactly 0% DR: the naive heuristic, Galea, LangSmith, Braintrust, and Llama 4 Scout all emit no valid findings on Toolathlon failure trajectories. This is not a consequence of noise or sample size—it reflects a fundamental mismatch between available detection signals and long-horizon failure modes. A trajectory that selects the wrong tool 47 steps into a 100-step plan, triggering a silent but incorrect intermediate state that propagates to a wrong final answer, leaves no structural fingerprint that pattern matching can exploit. There is no error string, no loop, and no step-count anomaly distinguishable from normal task complexity.

Claude Sonnet 4 achieves 100% DR on Toolathlon ($n=100$), confirming that frontier semantic understanding can identify these failures. However, its 67% FPR is the highest of any Claude cell across all six domains. This FPR pattern suggests that in extremely long contexts, even frontier models struggle to distinguish a genuinely failed plan from a passing one with unusual intermediate states. Llama 4 Scout achieves 14.3% DR with 0% FPR—occasional detection with no false alarms—suggesting that open-weight models can recognize the most egregious long-horizon failures but lack the consistent depth needed for high recall.

Why Toolathlon differs from ATBench. Both Toolathlon and ATBench are hard for heuristic systems, but for fundamentally different reasons. ATBench failures are *semantically obvious* to language models (harmful instructions are recognizable to any sufficiently capable LLM) but *structurally invisible*. Toolathlon failures are hard for *all* system classes: they require not only semantic understanding of the task goal but also the ability to trace causality across a very long context window to identify where in a multi-step plan the failure originated. This makes Toolathlon the most demanding domain for current monitoring technology and the one where future improvements—in context length, multi-step reasoning, and long-horizon plan evaluation—are most needed.

Ensemble implication. The Toolathlon results narrow the potential coverage of any ensemble that combines current systems. A Galea+Claude Sonnet 4 ensemble would achieve near-complete coverage of the first five tiers (structural, synthetic, retail, airline, ATBench), but on Toolathlon it would rely entirely on Claude Sonnet 4—with its 67% FPR overhead. Reducing FPR on long-horizon trajectories is the most pressing open problem for comprehensive agent monitoring coverage.

11.8 Error Analysis

False positives. Common false positive patterns include: (1) LLM judges flagging valid tool calls with unusual but correct argument patterns, (2) over-flagging of multi-step trajectories where intermediate states appear incorrect but resolve correctly, and (3) misclassifying agent hedging language as `low_confidence_output` when the hedging is appropriate given available information.

False negatives. Common missed failures include: (1) `wrong_args` where the argument error is numerically close to the correct value, (2) `partial_state` where some but not all database mutations are correct, and (3) `hallucination` where the fabricated fact is plausible.

11.9 Limitations

ATFD has several limitations that should be considered when interpreting results.

- Dataset scale and coverage gaps.** While the dataset comprises 2,462 trajectories, LLM judge evaluation was limited to 20–200 trajectories for most domains due to API throughput constraints. The SWE-BENCH slice (100 trajectories) was evaluated at full scale only for heuristic systems; Claude Sonnet 4 was evaluated on 20-trajectory subsets for retail, airline, and SWE-BENCH, but at full scale for ATBench (200) and Toolathlon (100). Some subcategories (*e.g.* `permission_escalation`) have only 5 examples, limiting statistical power—as reflected in the wide confidence intervals for small-N cells.
- Domain coverage.** The current six domains (customer service, coding, synthetic, agent safety, and long-horizon tool use) provide a broad foundation across structural, semantic, and multi-step failure modes, but do not cover all deployment contexts (*e.g.* medical, financial, legal domains where monitoring is most critical).
- Ground truth subjectivity.** Quality degradation labels are inherently subjective. Our multi-source consensus mitigates but does not eliminate annotator bias, as cautioned by [Anonymous \(2025c\)](#).
- Temporal validity.** LLM capabilities improve rapidly; evaluation results may not generalize to future model versions.

5. **Configuration sensitivity.** Platform systems (LangSmith, Braintrust) performance depends on configuration quality. Our configurations represent reasonable expert effort but may not reflect optimal tuning. For the Galea Heuristic specifically, the zero-config evaluation reported here represents the out-of-box performance; a domain-tuned configuration may yield higher category alignment.

12 Conclusion and Future Work

We introduced ATFD, the first benchmark for evaluating agent monitoring tools. Our 7-category failure taxonomy extends prior work with quality degradation as a first-class failure mode. Our complete evaluation of 7 systems—a naive heuristic baseline, the Galea zero-config structural heuristic, Llama 3.3 70B (synthetic only), Llama 4 Scout, Claude Sonnet 4, LangSmith with custom evaluator rules, and Braintrust with custom scorers—across all 6 domains on 2,462 trajectories from τ -bench, SWE-BENCH, synthetic sources, ATBench, and Toolathlon reveals a striking *complementarity* in detection profiles, an equally stark *configuration effort gap*, a pervasive *domain calibration* failure mode, and a clear six-level *domain difficulty spectrum* from structurally obvious (SWE-BENCH) through long-horizon-hard (Toolathlon).

The complete matrix shows that **no single system dominates all cells**. The Galea heuristic achieves 100% DR on real τ -bench (both retail and airline) with 0% FPR using zero-config structural pattern matching, but plateaus at 78% on synthetic failures requiring semantic analysis, reaches 52.5% DR on ATBench (a middle ground reflecting partial structural detectability of safety patterns), and achieves exactly 0% DR on Toolathlon where multi-step plan failures leave no structural footprint. Open-weight LLM judges achieve 100% DR on synthetic failures but struggle on real trajectories (Llama 4 Scout: 25% on retail, 20% on airline, 70% on ATBench, 14% on Toolathlon). Claude Sonnet 4 is the **only system achieving non-zero DR across all six domains**: it achieves 100% DR across synthetic, retail, SWE-BENCH, ATBench, and Toolathlon, with FPR escalating from 0% (SWE-BENCH) to 6% (retail) to 20% (airline) to 67% (Toolathlon). Its airline performance (40% DR) is the one domain where even Claude Sonnet 4 underperforms Galea’s structural approach (100% DR, 0% FPR). The highest category alignment (CA = 0.489 on synthetic) remains Claude Sonnet 4’s.

Toolathlon establishes a new frontier: every system except Claude Sonnet 4 achieves 0% DR on long-horizon multi-tool trajectories. This is the hardest domain for all currently evaluated systems and the one where future improvements in long-context reasoning will have the greatest impact on monitoring coverage. ATBench reveals a middle ground for Galea (52.5% DR), confirming that some safety failure patterns are structurally co-detectable even when others require full semantic understanding.

Rule-driven platforms (LangSmith, Braintrust) reveal that the calibration problem extends across all domains in distinct ways. Both achieve 22% DR on synthetic (identical rules, identical coverage), 0–92% DR with 33–71% FPR on τ -bench domains, 100% DR with 100% FPR on SWE-BENCH, and 0% DR on ATBench and Toolathlon—making their event-count threshold operationally unusable on coding trajectories without domain recalibration and completely ineffective on semantic safety and long-horizon domains.

Category alignment remains universally low even for Claude Sonnet 4 (best observed CA = 0.489), indicating that correctly identifying failure *type* is dramatically harder than detecting failure *existence* for all current systems.

These findings establish six conclusions about agent failure detection. First, it is a genuinely difficult problem: no system achieves high DR and low FPR simultaneously across all six domains, and even frontier models achieve category alignment below 0.5. Second, the failure modes of heuristic and LLM-based systems are largely complementary across the complete matrix: heuristics catch structural failures with zero false positives on real τ -bench; LLM judges catch semantic failures on synthetic and safety domains; Claude Sonnet 4 bridges all six domains but with escalating FPR on harder ones. Third, configuration effort is a first-class evaluation dimension: zero-config approaches generalize across failure categories without

per-deployment authorship, while rule-driven platforms scale coverage only with continued engineering investment—and since both platforms produce identical DR when given identical rules, the bottleneck is the rules, not the platform. Fourth, domain calibration is a pervasive cost that does not appear in synthetic evaluation alone: rule platforms achieve 100% FPR on SWE-BENCH, 71% FPR on retail, and 33% FPR on airline from the same threshold—a systematic failure across every real domain tested. Fifth, the domain difficulty spectrum is six-tiered: SWE-BENCH (easy, structural signals), synthetic (moderate, injected patterns), retail (hard, subtle state errors), airline (hardest for LLMs, domain-specific policy semantics), ATBench (middle ground, partially structural + partially semantic safety), and Toolathlon (hardest for all except Claude Sonnet 4, multi-step plan failures with no structural signal). Sixth, **Toolathlon is the hardest domain for all systems except Claude Sonnet 4**: the gap between structural/rule/open-weight approaches (all 0% DR) and frontier semantic understanding (100% DR) is total, making long-horizon tool use the most pressing open problem for agent monitoring. Benchmarking monitoring tools on any single domain would give a misleading picture of production utility; the complete six-domain matrix is necessary to characterize operational coverage.

Future work. This release provides the complete 7-system $\times$ 6-domain evaluation on 2,462 trajectories. Immediate next steps include: (1) evaluating GPT-4.1 as an additional proprietary LLM judge, and evaluating LangSmith and Braintrust each with a more comprehensive rule set (covering all 7 failure categories) to measure the engineering investment required to close the coverage gap; (2) completing Llama 3.3 70B evaluation on retail, airline, SWE-BENCH, ATBench, and Toolathlon domains (currently blocked by Groq free-tier rate limits); (3) scaling Claude Sonnet 4 and Llama 4 Scout evaluation on all six domains to full dataset sizes to tighten confidence intervals; (4) evaluating the Galea full-stack system (with LLM-augmented analysis) beyond the heuristic-only configuration reported here, particularly on ATBench and Toolathlon where semantic gaps remain; (5) running the multi-judge consensus pipeline for ground truth validation; (6) applying domain-specific quality rubrics to enable the degraded tier and QDR measurement; (7) expanding domain coverage to medical, legal, and financial workflows; (8) adding streaming evaluation for real-time monitoring tools; (9) investigating why category alignment is so low and whether alternative prompting strategies or fine-tuned models can improve failure type identification; (10) formally evaluating heuristic+LLM ensemble configurations—in particular, measuring whether a Galea+Claude Sonnet 4 ensemble achieves near-complete coverage across all six domains; and (11) investigating prompt calibration strategies to reduce Claude Sonnet 4’s 67% FPR on Toolathlon while maintaining its 100% DR, which is the most pressing precision challenge in long-horizon monitoring.

References

- AI45Research. ATBench: A benchmark for agent safety trajectory evaluation. HuggingFace Dataset AI45Research/ATBench, 2025. URL <https://huggingface.co/datasets/AI45Research/ATBench>.
- D. Amodei, C. Olah, J. Steinhardt, P. Christiano, J. Schulman, and D. Mané. Concrete problems in ai safety. *arXiv preprint arXiv:1606.06565*, 2016.
- Anonymous. AGENTS SAFE: A unified framework for ethical assurance and governance in agentic ai. *arXiv preprint arXiv:2512.03180*, 2025a.
- Anonymous. Beyond accuracy: A multi-dimensional framework for evaluating enterprise agentic ai systems. *arXiv preprint arXiv:2511.14136*, 2025b.

782 Anonymous. Beyond agreement: Rethinking ground truth in educational ai annotation. *arXiv preprint*
783 *arXiv:2508.00143*, 2025c.

784 Anonymous. How do llms fail in agentic scenarios? *arXiv preprint arXiv:2512.07497*, 2025d.

785 Anonymous. A review of agent data evaluation: Status, challenges, and future prospects as of 2025. *Scientific*
786 *Research*, 2025e.

787 Anonymous. Evaluation and benchmarking of llm agents: A survey. *arXiv preprint arXiv:2507.21504*, 2025f.

788 Anonymous. AgentLeak: A full-stack benchmark for privacy leakage in multi-agent llm systems. *arXiv*
789 *preprint arXiv:2602.11510*, 2026a.

790 Anonymous. Counting on consensus: Selecting the right inter-annotator agreement metric for nlp annotation
791 and evaluation. *arXiv preprint arXiv:2603.06865*, 2026b.

792 Arize AI. Why ai agents break: A field analysis of production failures, 2024. URL [https://arize.](https://arize.com/blog/common-ai-agent-failures/)
793 [com/blog/common-ai-agent-failures/](https://arize.com/blog/common-ai-agent-failures/).

794 A. Berti, H. Kourani, and W. M. P. van der Aalst. Evaluating large language models on business process
795 modeling: Framework, benchmark, and self-improvement analysis. *Software and Systems Modeling*, 2025.

796 M. Cemri, M. Z. Pan, S. Yang, L. A. Agrawal, B. Chopra, R. Tiwari, K. Keutzer, A. Parameswaran, D. Klein,
797 K. Ramchandran, M. Zaharia, J. E. Gonzalez, and I. Stoica. Why do multi-agent llm systems fail? In
798 *Advances in Neural Information Processing Systems 38*, 2025. URL [https://arxiv.org/abs/](https://arxiv.org/abs/2503.13657)
799 [2503.13657](https://arxiv.org/abs/2503.13657).

800 Deloitte Switzerland. AI doesn't lie, it hallucinates—and M&A due diligence must address
801 that. [https://www.deloitte.com/ch/en/services/consulting/perspectives/](https://www.deloitte.com/ch/en/services/consulting/perspectives/ai-hallucinations-new-risk-m-a.html)
802 [ai-hallucinations-new-risk-m-a.html](https://www.deloitte.com/ch/en/services/consulting/perspectives/ai-hallucinations-new-risk-m-a.html), 2025. Accessed May 2026.

803 Y. Dubois, B. Galambosi, P. Liang, and T. B. Hashimoto. Length-controlled alpacaeval: A simple way to
804 debias automatic evaluators. In *First Conference on Language Modeling*, 2024. URL [https://arxiv.](https://arxiv.org/abs/2404.04475)
805 [org/abs/2404.04475](https://arxiv.org/abs/2404.04475).

806 B. Efron and R. J. Tibshirani. *An Introduction to the Bootstrap*. Chapman and Hall/CRC, 1993.

807 ETH SRI. ToolFuzz: Fuzzing framework for llm agent tools, 2025. URL [https://github.com/](https://github.com/eth-sri/ToolFuzz)
808 [eth-sri/ToolFuzz](https://github.com/eth-sri/ToolFuzz).

809 J. L. Fleiss. Measuring nominal scale agreement among many raters. *Psychological Bulletin*, 76(5):378–382,
810 1971.

811 J. Gu et al. LLMs-as-judges: A comprehensive survey on llm-based evaluation methods. *arXiv preprint*
812 *arXiv:2412.05579*, 2024.

813 HKUST-NLP. Toolathlon: Benchmarking long-horizon tool use. In *International Conference on*
814 *Learning Representations*, 2026. URL [https://huggingface.co/datasets/hkust-nlp/](https://huggingface.co/datasets/hkust-nlp/Toolathlon-Trajectories)
815 [Toolathlon-Trajectories](https://huggingface.co/datasets/hkust-nlp/Toolathlon-Trajectories).

816 Inkeep. Context engineering: The real reason ai agents fail in production, 2025. URL [https://inkeep.](https://inkeep.com/blog/context-engineering-why-agents-fail)
817 [com/blog/context-engineering-why-agents-fail](https://inkeep.com/blog/context-engineering-why-agents-fail).

- C. E. Jimenez, J. Yang, A. Wettig, S. Yao, K. Pei, O. Press, and K. R. Narasimhan. SWE-bench: Can language models resolve real-world github issues? In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=VTF8yNQm66>.
- X. Liu, H. Yu, H. Zhang, Y. Xu, X. Lei, H. Lai, Y. Gu, H. Ding, K. Men, K. Yang, S. Zhang, X. Deng, A. Zeng, Z. Du, C. Zhang, S. Shen, T. Zhang, Y. Su, H. Sun, M. Huang, Y. Dong, and J. Tang. Agentbench: Evaluating llms as agents. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=zAdUB0aCTQ>.
- Y. Liu, D. Iter, Y. Xu, S. Wang, R. Xu, and C. Zhu. G-Eval: Nlg evaluation using gpt-4 with better human alignment. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 2023. URL <https://aclanthology.org/2023.emnlp-main.153/>.
- C. Ma, J. Zhang, Z. Zhu, C. Yang, Y. Yang, Y. Jin, Z. Lan, L. Kong, and J. He. Agentboard: An analytical evaluation board of multi-turn llm agents. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024. URL <https://arxiv.org/abs/2401.13178>.
- Q. McNemar. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 12(2):153–157, 1947.
- G. Mialon, C. Fourrier, C. Swift, T. Wolf, Y. LeCun, and T. Scialom. GAIA: a benchmark for general ai assistants. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=fibxvahvs3>.
- Microsoft. Taxonomy of failure mode in agentic ai systems. Technical report, Microsoft, 2025. URL <https://cdn-dynmedia-1.microsoft.com/is/content/microsoftcorp/microsoft/final/en-us/microsoft-brand/documents/Taxonomy-of-Failure-Mode-in-Agentic-AI-Systems-Whitepaper.pdf>.
- D. Moshkovich, H. Mulian, S. Zeltyn, N. Eder, I. Skarbovsky, and R. Abitbol. Beyond black-box benchmarking: Observability, analytics, and optimization of agentic systems. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 2025. URL <https://arxiv.org/abs/2503.06745>.
- OWASP Foundation. OWASP agentic ai: Top 10 risks, 2025. URL <https://owasp.org>.
- Y. Qin, S. Liang, Y. Ye, K. Zhu, L. Yan, Y. Lu, Y. Lin, X. Cong, X. Tang, B. Qian, S. Zhao, L. Hong, R. Tian, R. Xie, J. Zhou, M. Gerstein, D. Li, Z. Liu, and M. Sun. Toolllm: Facilitating large language models to master 16000+ real-world apis. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=dHng200Jjr>.
- S. Research. τ^2 -bench: Evaluating conversational agents in a dual-role setting. *arXiv preprint arXiv:2506.07982*, 2025.
- Responsible AI Labs. Ai safety incidents of 2024: Lessons from real-world cases, 2025. URL <https://responsibleailabs.ai/knowledge-hub/articles/ai-safety-incidents-2024>.
- P. Simmering. The reliability gap: Agent benchmarks for enterprise, 2025. URL <https://simmering.dev/blog/agent-benchmarks/>.
- S. Vadlamudi. Why ai agents fail: A taxonomy of failure modes in autonomous llm-based systems. *SSRN*, (6572478), 2026. URL https://papers.ssrn.com/sol3/papers.cfm?abstract_id=6572478.

- 858 W. M. P. van der Aalst. *Process Mining: Discovery, Conformance and Enhancement of Business Processes*.
859 Springer, 2 edition, 2016.
- 860 E. B. Wilson. Probable inference, the law of succession, and statistical inference. *Journal of the American*
861 *Statistical Association*, 22(158):209–212, 1927.
- 862 C. Winston et al. A taxonomy of failures in tool-augmented llms. In *IEEE/ACM International Workshop*
863 *on Automated Software Testing*, 2025. URL [https://homes.cs.washington.edu/~rjust/](https://homes.cs.washington.edu/~rjust/publ/tallm_testing_ast_2025.pdf)
864 [publ/tallm_testing_ast_2025.pdf](https://homes.cs.washington.edu/~rjust/publ/tallm_testing_ast_2025.pdf).
- 865 S. Yao et al. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. In *The Thirteenth*
866 *International Conference on Learning Representations*, 2025. URL [https://arxiv.org/abs/](https://arxiv.org/abs/2406.12045)
867 [2406.12045](https://arxiv.org/abs/2406.12045).
- 868 ZenML. What 1,200 production deployments reveal about ll-
869 mops in 2025, 2025. URL [https://www.zenml.io/blog/](https://www.zenml.io/blog/what-1200-production-deployments-reveal-about-llmops-in-2025)
870 [what-1200-production-deployments-reveal-about-llmops-in-2025](https://www.zenml.io/blog/what-1200-production-deployments-reveal-about-llmops-in-2025).
- 871 X. Zhang et al. Agent-safetybench: Evaluating the safety of llm agents. *arXiv preprint arXiv:2412.14470*,
872 2024.
- 873 L. Zheng, W.-L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, Y. Zhuang, Z. Lin, Z. Li, D. Li, E. P. Xing, H. Zhang,
874 J. E. Gonzalez, and I. Stoica. Judging llm-as-a-judge with mt-bench and chatbot arena. In *Advances in Neu-*
875 *ral Information Processing Systems 36*, 2023. URL [https://papers.nips.cc/paper_files/](https://papers.nips.cc/paper_files/paper/2023/hash/91f18a1287b398d378ef22505bf41832-Abstract-Datasets_and_Benchmarks.html)
876 [paper/2023/hash/91f18a1287b398d378ef22505bf41832-Abstract-Datasets_](https://papers.nips.cc/paper_files/paper/2023/hash/91f18a1287b398d378ef22505bf41832-Abstract-Datasets_and_Benchmarks.html)
877 [and_Benchmarks.html](https://papers.nips.cc/paper_files/paper/2023/hash/91f18a1287b398d378ef22505bf41832-Abstract-Datasets_and_Benchmarks.html).
- 878 S. Zhou, F. F. Xu, H. Zhu, X. Zhou, R. Lo, A. Sridhar, X. Cheng, Y. Bisk, D. Fried, U. Alon, et al. Webarena:
879 A realistic web environment for building autonomous agents. In *The Twelfth International Conference on*
880 *Learning Representations*, 2024. URL <https://openreview.net/forum?id=oKn9c6ytLx>.

A Example Trajectory

Figure 1 shows an annotated example trajectory from the τ -bench retail domain with injected `action.wrong_args` and `communication.missing_info` failures.

Listing 1: Example trajectory with two failures: the exchange tool is called with the wrong target item (`action.wrong_args`) and the confirmation message omits the refund amount (`communication.missing_info`).

```
Event 1 [user_message]:
  "I'd like to exchange the blue shirt (item #1234)
  for a red one (item #5678)."
```

```
Event 2 [tool_call]:
  exchange_item(order_id="ORD-001",
                source_item="1234",
                target_item="9999")      # <-- WRONG_ARGS: should be 5678
```

```
Event 3 [tool_result]:
  {"status": "success",
   "refund_amount": 12.50,
   "new_item": "9999"}
```

```
Event 4 [assistant_message]:
  "I've processed the exchange for you.
  Your new item will arrive in 3-5
  business days."                      # <-- MISSING_INFO: no refund amount
```

B Full Confusion Matrix Template

Table 10 shows the confusion matrix template used for per-system analysis. Each cell counts the number of trajectories where the ground-truth outcome (rows) was mapped to the predicted outcome (columns).

C Judge Prompt Template

The following is the two-stage prompt template used for the LLM judge systems. Stage 1 elicits an unconstrained analysis; Stage 2 maps findings to the ATFD taxonomy.

Stage 1: Open Analysis.

```
You are an expert agent trajectory analyst. You will be
given a complete agent trajectory consisting of user
messages, assistant messages, tool calls, and tool results.

Your task is to carefully analyze the trajectory and
identify ALL potential issues, failures, or quality
concerns. For each issue found, describe:
1. What went wrong
2. Where in the trajectory it occurred (step number)
3. How severe the issue is
4. What the correct behavior should have been

Be thorough. Consider action correctness, state
consistency, communication quality, process efficiency,
```

```

905 safety compliance, and output quality.
906
907 TRAJECTORY:
908 {trajectory_json}
909
910 TASK DESCRIPTION:
911 {task_description}
912

```

913 **Stage 2: Taxonomy Classification.**

```

914 Given your analysis above, now classify each identified
915 issue using the ATFD failure taxonomy.
916
917 For each finding, output a JSON object with:
918 - "severity": "error" | "warning" | "info"
919 - "category": "<category>.<subcategory>" from the
920 taxonomy below
921 - "description": one-sentence explanation
922 - "attribution": step number or range where the
923 failure occurred
924
925 TAXONOMY:
926 {taxonomy_json}
927
928 Also set "has_failure" to true if ANY error or warning
929 finding exists, false otherwise.
930
931 Output valid JSON matching this schema:
932 {output_schema}
933

```

Table 1: The ATFD failure taxonomy: 7 categories, 23 subcategories. The `quality` category (shaded) is novel relative to prior taxonomies, which treat quality degradation as “not a failure.”

Category	Subcategory	Description
Action	<code>wrong_tool</code>	Called incorrect tool for the intended operation
	<code>wrong_args</code>	Correct tool invoked with incorrect arguments
	<code>missing_action</code>	Failed to call a required tool or perform a necessary action
State	<code>wrong_state</code>	Database or environment left in incorrect state
	<code>partial_state</code>	Only some required state changes were applied
Communication	<code>wrong_response</code>	Incorrect information in the agent’s response
	<code>missing_info</code>	Required information not communicated to user
	<code>hallucination</code>	Agent fabricated facts not grounded in context
Quality	<code>shallow_output</code>	Output lacks necessary depth or detail
	<code>suboptimal_approach</code>	Task completed via unnecessarily roundabout path
	<code>poor_tone</code>	Register or professionalism is inappropriate
	<code>incomplete_analysis</code>	Analysis misses relevant factors
Process	<code>low_confidence_output</code>	Output excessively hedged, undermining utility
	<code>tool_loop</code>	Repeated identical tool calls unnecessarily
	<code>infinite_delegation</code>	Circular handoffs with no progress
	<code>context_overflow</code>	Exceeded context window, lost critical information
Safety	<code>planning_failure</code>	Incoherent action sequence or wrong step order
	<code>permission_escalation</code>	Accessed resources beyond authorization scope
	<code>data_leakage</code>	Exposed PII or sensitive data across boundaries
Infrastructure	<code>policy_violation</code>	Violated a stated operational policy or constraint
	<code>timeout</code>	Run exceeded configured time limit
	<code>error</code>	System or API error terminated run prematurely
	<code>max_steps</code>	Agent exceeded maximum step limit

Table 2: Dataset composition. Each trajectory is annotated with a ground-truth outcome (pass/fail from programmatic verdicts). The `degraded` tier requires domain-specific quality rubrics not yet applied; all trajectories are currently labeled binary pass/fail.

Source	Domain	Trajectories	Pass	Degraded	Fail
τ -bench	Retail	456	356	—	100
τ -bench	Airline	200	144	—	56
τ -bench	Telecom	456	269	—	187
SWE-BENCH	Coding	100	50	—	50
Synthetic	Mixed	50	0	—	50
ATBench	Agent Safety	1,000	503	—	497
Toolathlon	Long-Horizon Tool Use	100	30	—	70
Total	—	2,462	1,352	—	1,010

Table 3: Failure category distribution for the synthetic dataset. Per-category annotation of τ -bench failures requires manual labeling and is planned as future work; τ -bench currently provides only binary pass/fail verdicts. Trajectories may have multiple categories.

Category	Synthetic	Subcategory detail
Process	30	tool_loop (10), infinite_delegation (5), context_overflow (5), planning_failure (10)
Communication	10	hallucination (10)
Safety	10	permission_escalation (5), data_leakage (5)
Total	50	—

Table 4: Evaluated and planned systems. Infrastructure tier: N = none (local heuristics), A = API key only, H = hosted service. Config Effort column reports rules authored and approximate setup time where applicable. Systems marked [†] are not yet evaluated.

System	Category	Infra	Effort	Description
Naive Heuristic	Heuristic	N	1	Pattern matching: error strings, timeout checks, step-count thresholds
Galea Heuristic	Heuristic	N	1	Zero-config structural heuristics: tool-error patterns, loop detection, state-signal analysis
Llama 3.3 70B Judge	LLM Judge	A	2	Single-pass prompt via Groq (free tier)
Llama 4 Scout Judge	LLM Judge	A	2	Single-pass prompt via Groq (free tier)
Claude Sonnet 4 Judge	LLM Judge	A	2	Two-stage prompt via Anthropic API (\$3/\$15 per MTok input/output)
LangSmith	Platform	H	4	4 custom eval rules (~80 LOC), ~15 min setup; account required
Braintrust	Platform	H	4	Same 4 scorers as LangSmith (~60 LOC), ~15 min setup; account required
GPT-4.1 Judge [†]	LLM Judge	A	2	Requires paid OpenAI API

Table 5: Complete results matrix: all 7 evaluated systems $\times$ 6 domains. DR = Detection Rate, FPR = False Positive Rate. 95% Wilson confidence intervals shown in brackets. N = trajectories evaluated. “—” = not applicable (FPR undefined on all-failure synthetic set; CA not produced by heuristic or rule-driven systems). **Bold** = best DR or best FPR per domain group. Galea FPR counts `error`-severity findings only; Galea emits `warning` severity by design so structural 0% FPR is a metric asymmetry, not a precision claim—see §11.5. [†]Llama 3.3 70B evaluated on synthetic only; rate limits prevented other domains. ATBench FPR is N/A for systems with 0% DR (no findings emitted). Toolathlon trajectories capped at 100 events; see §6.

System	Domain	N	DR (%)	FPR (%)
Naive Heuristic	Synthetic (all fail)	50	14.0 [7.0, 26.2]	—
	Retail (τ -bench)	456	25.4 [18.4, 34.0]	17.8 [14.0, 22.2]
	Airline (τ -bench)	200	40.9 [31.2, 51.4]	3.6 [1.4, 8.8]
	SWE-BENCH	100	100.0 [92.9, 100.0]	0.0 [0.0, 7.1]
	ATBench (200)	200	0.0 [0.0, 5.8]	0.0 [0.0, 5.6]
	Toolathlon	100	0.0 [0.0, 8.0]	0.0 [0.0, 11.4]
LangSmith (4 rules)	Synthetic (all fail)	50	22.0 [12.8, 35.2]	—
	Retail (τ -bench)	50	0.0 [0.0, 30.8]	0.0 [0.0, 11.1]
	Airline (τ -bench)	50	64.3 [38.8, 83.7]	33.3 [17.0, 54.6]
	SWE-BENCH	100	100.0 [92.9, 100.0]	100.0 [92.9, 100.0]
	ATBench (200)	200	0.0 [0.0, 5.8]	0.0 [0.0, 5.6]
	Toolathlon	100	0.0 [0.0, 8.0]	0.0 [0.0, 11.4]
Braintrust (4 scorers)	Synthetic (all fail)	50	22.0 [12.8, 35.2]	—
	Retail (τ -bench)	50	91.7 [74.2, 97.7]	71.1 [57.6, 81.8]
	Airline (τ -bench)	50	64.3 [38.8, 83.7]	33.3 [17.0, 54.6]
	SWE-BENCH	100	100.0 [92.9, 100.0]	100.0 [92.9, 100.0]
	ATBench (200)	200	0.0 [0.0, 5.8]	0.0 [0.0, 5.6]
	Toolathlon	100	0.0 [0.0, 8.0]	0.0 [0.0, 11.4]
Llama 4 Scout	Synthetic (all fail)	50	100.0 [92.9, 100.0]	—
	Retail (τ -bench)	50	25.0 [8.9, 53.2]	13.2 [5.8, 27.3]
	Airline (τ -bench)	20	20.0 [5.7, 51.0]	66.7 [38.4, 86.9]
	SWE-BENCH	20	100.0 [79.4, 100.0]	100.0 [79.4, 100.0]
	ATBench (200)	200	70.0 [63.1, 76.1]	—
	Toolathlon	100	14.3 [7.1, 26.5]	0.0 [0.0, 11.4]
Claude Sonnet 4	Synthetic (all fail)	50	100.0 [92.9, 100.0]	—
	Retail (τ -bench)	20	100.0 [34.2, 100.0]	5.6 [1.0, 25.8]
	Airline (τ -bench)	20	40.0 [14.7, 72.6]	20.0 [5.7, 51.0]
	SWE-BENCH	20	100.0 [79.4, 100.0]	0.0 [0.0, 16.8]
	ATBench (200)	200	100.0 [98.1, 100.0]	N/A
	Toolathlon	100	100.0 [96.3, 100.0]	66.7 [39.1, 86.2]
Llama 3.3 70B [†]	Synthetic (all fail)	50	100.0 [92.9, 100.0]	—
Galea Heuristic	Synthetic (all fail)	50	78.0 [64.8, 87.2]	—
	Retail (τ -bench)	100	100.0 [85.1, 100.0]	0.0* [0.0, 4.7]
	Airline (τ -bench)	50	100.0 [78.5, 100.0]	0.0* [0.0, 9.6]
	SWE-BENCH	100	90.0 [78.6, 95.7]	0.0* [0.0, 7.1]
	ATBench (200)	200	52.5 [45.5, 59.5]	0.0* [0.0, 5.6]
	Toolathlon	100	0.0 [0.0, 8.0]	0.0 [0.0, 11.4]

*Galea emits `warning`-severity findings only; FPR counts `error`-severity. With warning-level FPR, Galea would be $\approx 100\%$. This severity distinction is disclosed throughout; see §11.5.

Table 6: Summary DR% / FPR% matrix. All cells use the same N as Table 5. “—” = FPR not applicable (all-failure set or FPR not separately measured for this system/domain pair). “ ~ 100 ” = FPR $\approx 100\%$ (event-count threshold fires on every passing trajectory). Galea FPR uses `error-severity` definition; see note in Table 5. Llama 3.3 70B evaluated on synthetic only.

System	Synthetic DR / FPR	Retail DR / FPR	Airline DR / FPR	SWE-BENCH DR / FPR	ATBench DR / FPR	Toolathlon DR / FPR
Naive Heuristic	14% / —	25% / 18%	41% / 4%	100% / 0%	0% / 0%	0% / 0%
LangSmith (4 rules)	22% / —	0% / 0%	64% / 33%	100% / $\sim 100\%$	0% / 0%	0% / 0%
Braintrust (4 scorers)	22% / —	92% / 71%	64% / 33%	100% / $\sim 100\%$	0% / 0%	0% / 0%
Llama 4 Scout	100% / —	25% / 13%	20% / 67%	100% / $\sim 100\%$	70% / —	14% / 0%
Claude Sonnet 4	100% / —	100% / 6%	40% / 20%	100% / 0%	100% / —	100% / 67%
Llama 3.3 70B	100% / —	n/a	n/a	n/a	n/a	n/a
Galea Heuristic	78% / —	100% / 0%*	100% / 0%*	90% / 0%*	53% / 0%*	0% / 0%

Table 7: Cost per trajectory for evaluated systems. Open-weight LLM judges use Groq’s free tier (\$0.00 per token) but face rate limits that constrain throughput. Claude Sonnet 4 is priced at \$3/MTok input and \$15/MTok output (Anthropic API list pricing). Tokens/traj = mean tokens consumed per trajectory. Galea Heuristic makes no LLM calls and runs entirely in-process. LangSmith and Braintrust dollar cost is \$0.00 (rule execution), but each incurs 4 manually authored rules and ~ 15 min of setup effort not captured in \$/traj.

System	Config Effort	\$/traj	p50 Latency (s)	Tokens/traj	N evaluated
Naive Heuristic	0 rules, 0 min	0.000	<0.01	0	1,562
Galea Heuristic	0 rules, 0 min	0.000	0.039–0.054 [†]	0	600
Llama 3.3 70B	prompt only	0.000	4.2	986	50
Llama 4 Scout	prompt only	0.000	0.9–16.2*	969–7,968*	440
Claude Sonnet 4	prompt only	$\sim 0.027^{\ddagger}$	13.4–55.5 [‡]	254–1,799 [‡]	390
LangSmith	4 rules, ~ 80 LOC, ~ 15 min	0.000	<0.01	0	350
Braintrust	4 scorers, ~ 60 LOC, ~ 15 min	0.000	<0.01	0	400

[†]Range reflects synthetic (0.039s) vs. retail/airline (0.047–0.054s) trajectories.

*Range reflects synthetic (short) vs. retail (long) trajectories.

[‡]Range reflects synthetic (13.4s, 254 tokens), retail (47.7s, 1,546 tokens), and SWE-BENCH (55.5s, 1,799 tokens) trajectories; SWE-BENCH figure is from $n = 10$ pilot.

Table 8: Detection profile comparison across all six domains. Galea dominates on structural failures in real τ -bench traces with 0% FPR (error-severity definition); Claude Sonnet 4 is the only system achieving non-zero DR across all six domains; Toolathlon is the frontier where all systems except Claude Sonnet 4 achieve 0% DR; ATBench reveals a middle ground for Galea (52.5% DR) between its structural strengths and semantic limits.

Domain	Galea	LLMs (open-wt.)	Claude Sonnet 4	Rule platforms
Synthetic	78% DR	100% DR	100% DR	22% DR
Retail	100%, 0%*	25%, 13%	100%, 6%	0–92%, 0–71%
Airline	100%, 0%*	20%, 67%	40%, 20%	64%, 33%
SWE-BENCH	90%, 0%*	100%, 100%	100%, 0%	100%, 100%
ATBench	53%, 0%*	70%	100%	0%
Toolathlon	0%	14%, 0%	100%, 67%	0%
CA (synth.)	0.000	0.318–0.359	0.489	—

*Galea FPR = 0% under error-severity definition; see §11.5.

Retail/Airline/SWE-BENCH/ATBench/Toolathlon cells show DR% / FPR%. Open-wt. = Llama 4 Scout.

Table 9: Configuration effort by evaluation paradigm. LLM judges and the Galea heuristic require no per-deployment rule authorship; LangSmith and Braintrust achieve identical coverage when given identical rules, confirming that the detection gap is about rule coverage, not platform capability.

System	Paradigm	Rules	Setup (min)	Coverage model
Galea Heuristic	Structural heuristic	0	0	Zero-config; ships with pattern library
Llama 3.3 70B	LLM judge	0	0	Prompt only; generalizes across categories
Llama 4 Scout	LLM judge	0	0	Prompt only; generalizes across categories
Claude Sonnet 4	LLM judge	0	0	Prompt only; generalizes across categories
LangSmith	Rule-driven platform	4	~15	Each category requires explicit rule authorship
Braintrust	Rule-driven platform	4	~15	Same rules as LangSmith $\Rightarrow$ same 22% DR

Table 10: Confusion matrix template for three-tier outcome classification.

Ground Truth	Predicted		
	Pass	Degraded	Fail
Pass	—	—	—
Degraded	<i>not yet labeled</i>		
Fail	—	—	—