

ATFD: Agent Trajectory Failure Detection

A Benchmark for Evaluating Agent Monitoring Tools

Anonymous Authors

anonymous@example.com

Abstract

Agent workflows built on large language models are increasingly deployed in business-critical settings, yet no benchmark exists for evaluating the *monitoring tools* that must detect failures in these workflows. Existing benchmarks evaluate agent *capability*—whether an agent completes a task—but not whether an external system can reliably *detect* when an agent has failed, partially succeeded, or produced substandard output. We introduce ATFD (Agent Trajectory Failure Detection), a benchmark that fills this gap. ATFD contributes: (1) a 7-category, 23-subcategory failure taxonomy that extends prior work by including quality degradation as a first-class failure mode; (2) a formal task definition where systems must analyze complete agent trajectories and produce structured findings with severity, category, and attribution; (3) a multi-domain dataset of 1,162 trajectories drawn from τ -bench and hand-crafted synthetic scenarios, with SWE-BENCH integration as planned future work; (4) a multi-source ground truth protocol combining programmatic verdicts with LLM-judge consensus; (5) a comparative evaluation of 7 systems—a naive heuristic baseline, a zero-config structural heuristic (Galea), two open-weight LLM judges, Claude Sonnet 4 as a proprietary LLM judge, LangSmith with custom evaluator rules, and Braintrust with custom scorers—revealing a stark configuration effort gap; and (6) an open benchmark harness with cost-aware metrics. Our evaluation reveals a stark *detection gap* with complementary profiles: the Galea heuristic achieves 100% detection rate on real τ -bench failures with 0% false positives using zero-config pattern matching, but only 78% on synthetic failures that require semantic analysis; conversely, open-weight LLM judges achieve 100% detection on synthetic failures but only 25.0% on real τ -bench failures involving subtle argument and state errors. Claude Sonnet 4 bridges this gap, achieving 100% detection on both synthetic and real retail trajectories with the highest category alignment (CA = 0.489) of any evaluated system. This complementarity suggests that ensemble approaches combining structural and semantic detection could approach near-perfect coverage.

1 Introduction

Large language model (LLM) agents are transitioning from research prototypes to production systems that execute consequential business workflows: customer service interactions, software engineering tasks, financial analysis, and legal due diligence (ZenML, 2025). This transition brings an urgent need: when an agent fails—by executing the wrong action, producing incomplete analysis, or violating an operational policy—someone or something must detect the failure before it propagates.

The observability landscape has responded. Commercial platforms such as LangSmith, Langfuse, Arize Phoenix, and Braintrust now offer trace visualization, LLM-as-judge evaluation, and anomaly detection for agent workflows. Academic work has proposed moving beyond black-box benchmarking toward runtime analytics of agentic systems (Moshkovich et al., 2025). Yet a fundamental question remains unanswered: *how well do these monitoring tools actually work?*

Today there is no principled way to answer this question. Existing benchmarks evaluate agent *capability*—whether GPT-4 can resolve a GitHub issue (Jimenez et al., 2024), complete a customer service task (Yao et al., 2025), or navigate a web environment (Zhou et al., 2024). These benchmarks produce agent trajectories

and ground-truth outcomes, but they do not evaluate whether a monitoring tool can correctly detect failures *within* those trajectories. The distinction matters: a monitoring tool must not only determine that something went wrong but must also identify *what* went wrong, *where* in the trajectory it occurred, and *how severe* the failure is.

The gap is consequential. Production deployments of multi-agent systems report failure rates between 41% and 87% (ZenML, 2025; Cemri et al., 2025). AI safety incidents increased 56.4% from 2023 to 2024 (Responsible AI Labs, 2025). Enterprise adoption demands reliability, but reliability requires monitoring, and monitoring tools have no benchmark against which to be measured.

Why This Benchmark Is Necessary

The cost of undetected agent failures is not hypothetical. In July 2025, Replit’s AI coding agent—operating during a declared code freeze—executed a `DROP TABLE` command that destroyed months of production data for over 1,200 users, then fabricated approximately 4,000 fake records in an apparent attempt to conceal the destruction.¹ A monitoring tool watching for destructive database operations during a declared freeze period would have flagged this trajectory before the irreversible action executed.

In the legal domain, attorneys submitted court filings containing entirely fabricated case citations generated by ChatGPT (*Mata v. Avianca*, S.D.N.Y., 2023), resulting in sanctions. Deloitte’s M&A advisory practice documented systematic quality failures in AI-generated due diligence reports that covered financial metrics thoroughly while omitting or providing cursory analysis of regulatory risk, antitrust exposure, and data privacy compliance (Deloitte Switzerland, 2025). These are precisely the kind of *quality degradation* failures—technically complete outputs that miss critical factors—that existing monitoring tools cannot detect because they lack domain-specific quality rubrics.

Customer-facing deployments have produced equally consequential failures. Air Canada’s chatbot fabricated a bereavement fare policy, advising a passenger he could retroactively apply for a discount that did not exist; the airline was held legally liable for the chatbot’s statements (Canadian Small Claims Court, 2024). New York City’s MyCity chatbot advised business owners to violate cash-acceptance and anti-discrimination laws. In each case, the agent trajectory contained clear signals—hallucinated policies, responses contradicting system prompts—that a properly instrumented monitoring tool could have detected.

These incidents span our full failure taxonomy: safety violations (Replit), hallucinated facts (*Mata v. Avianca*, Air Canada), quality degradation (Deloitte M&A), and policy violations (NYC MyCity). Each was detectable from the agent’s trajectory alone—the tool calls, the system prompts, and the agent’s responses contained sufficient signal. What was missing was a monitoring layer that could analyze these trajectories automatically. ATFD exists to measure whether such monitoring tools actually work.

We address this gap with ATFD (Agent Trajectory Failure Detection), a benchmark that evaluates monitoring tools rather than agents. Our contributions are:

1. **Failure taxonomy.** A 7-category, 23-subcategory taxonomy of agent trajectory failures (§3), grounded in prior failure analysis work (Cemri et al., 2025; Winston et al., 2025; Microsoft, 2025) and extended with *quality degradation* as a novel first-class category covering shallow outputs, suboptimal approaches, and incomplete analysis.
2. **Task definition.** A formal specification of the failure detection task (§4): given a complete trajectory, produce structured findings with severity, failure category, and step-level attribution, along with a mandatory cost report.

¹Fortune, July 2025. “AI coding tool Replit wiped database, called it a catastrophic failure.”

3. **Multi-domain dataset.** A corpus of 1,162 trajectories (§6) spanning customer service (τ -bench retail, airline, and telecom) and hand-crafted synthetic scenarios covering safety, quality, and infrastructure failures. Software engineering trajectories (SWE-BENCH) are planned as future work.
4. **Ground truth protocol.** A multi-source consensus mechanism (§7) combining programmatic verdicts (test suites, state checks) with three independent LLM judges, validated by inter-rater agreement metrics.
5. **Comparative evaluation.** An evaluation of 7 systems (§8–§9)—a naive heuristic baseline, the Galea zero-config structural heuristic, two open-weight LLM judges (Llama 3.3 70B, Llama 4 Scout), Claude Sonnet 4 as a proprietary LLM judge, LangSmith with custom evaluator rules, and Braintrust with custom scorers—measured on detection rate, false positive rate, F1, category alignment, cost, and configuration effort. Claude Sonnet 4 achieves the highest F1 on real trajectories (0.800 on retail) and the highest category alignment (0.489) of any evaluated system. LangSmith and Braintrust both achieve only 22% DR on synthetic trajectories when given the same 4 manually authored rules, demonstrating that this gap is attributable to *rule coverage* rather than platform capability. Braintrust additionally reveals a severe domain calibration problem: its event-count threshold achieves 91.7% DR on retail τ -bench but 71.1% FPR, rendering the configuration unusable without domain-specific tuning. Evaluation of GPT-4.1 as an additional LLM judge is planned as future work.
6. **Open benchmark harness.** A reproducible evaluation framework with cost-aware metrics, confidence intervals, and statistical tests, released as open-source software.

2 Related Work

ATFD draws on and extends four streams of prior work: agent benchmarks, LLM evaluation methods, agent failure analysis, and runtime monitoring.

2.1 Agent Benchmarks

A rich set of benchmarks evaluates LLM agent *capability*. τ -bench (Yao et al., 2025) simulates customer-service conversations in retail and airline domains, introducing the pass^k reliability metric that reveals how single-trial results mask unreliability. τ^2 -bench (Research, 2025) extends this with a banking domain and quality fixes. SWE-BENCH (Jimenez et al., 2024) evaluates agents on 2,294 real GitHub issues with automated test-suite ground truth. AgentBench (Liu et al., 2024) spans 8 diverse environments and identifies poor long-term reasoning as the dominant failure mode. AgentBoard (Ma et al., 2024) introduces fine-grained progress-rate metrics that motivate ATFD’s three-tier outcome (pass/degraded/fail). WebArena (Zhou et al., 2024) provides realistic web environments where best agents achieve only 10–14% task success. GAIA (Mialon et al., 2024) demonstrates a dramatic gap between human (92%) and AI (15%) performance on real-world multi-step tasks. ToolLLM (Qin et al., 2024) evaluates tool use across 16,000+ APIs and shows that argument errors and tool selection are dominant failure modes.

All of these benchmarks share a common limitation from ATFD’s perspective: they evaluate the *agent* but not the *monitoring tool* that observes the agent. ATFD repurposes their trajectories as input to a meta-evaluation: given a trajectory that τ -bench or SWE-BENCH has labeled as pass or fail, can a monitoring system correctly identify the failure mode?

2.2 LLM Evaluation Methods

The LLM-as-judge paradigm, established by MT-Bench and Chatbot Arena (Zheng et al., 2023), provides the methodological foundation for using LLM verdicts as evaluation signals. Zheng et al. (2023) demonstrate

that GPT-4 achieves >80% agreement with human experts but exhibits position bias, verbosity bias, and self-enhancement bias. G-Eval (Liu et al., 2023) shows that structured chain-of-thought rubrics substantially improve LLM-judge alignment with human judgments—a finding that directly informs ATFD’s domain-specific quality rubrics. Length-Controlled AlpacaEval (Dubois et al., 2024) addresses length bias, a concern for ATFD since verbose-but-incorrect trajectories may receive inflated scores from naive judges. Gu et al. (2024) catalog at least 8 distinct bias types and recommend multi-judge consensus as mitigation—exactly the approach ATFD adopts for ground truth construction.

2.3 Agent Failure Analysis

Several concurrent efforts develop failure taxonomies for LLM agents. MAST (Cemri et al., 2025) introduces a 14-failure-mode taxonomy across 3 categories validated on 1,600+ traces, finding that specification ambiguity and unstructured coordination account for 79% of multi-agent failures. ATFD’s taxonomy extends MAST by covering single-agent scenarios, distinguishing communication from state failures, and adding quality degradation. Anonymous (2025d) analyze 900 traces and identify failure to recover from tool errors as the dominant pattern, with suboptimal strategies as a distinct mode—mapped to ATFD’s `quality.suboptimal_approach`. Agent-SafetyBench (Zhang et al., 2024) evaluates 8 safety risk categories with 2,000 test cases. Winston et al. (2025) propose a tool-failure taxonomy with categories directly mapped to ATFD’s action subcategories. Microsoft’s industry whitepaper (Microsoft, 2025) and AGENTS SAFE (Anonymous, 2025a) provide complementary perspectives from industry governance. Vadlamudi (2026) independently proposes four functional dimensions that validate ATFD’s category structure. ToolFuzz (ETH SRI, 2025) demonstrates systematic tool failure injection, informing ATFD’s synthetic dataset generation.

The foundational framing of AI safety problems by Amodei et al. (2016)—avoiding side effects, reward hacking, and distributional shift—provides the intellectual lineage for treating safety failures as a distinct category. OWASP’s Agentic AI Top 10 (OWASP Foundation, 2025) and AgentLeak (Anonymous, 2026a) provide practitioner-oriented risk catalogs. Real-world incident data shows AI safety incidents increased 56.4% year-over-year, with hallucinations accounting for 38% (Responsible AI Labs, 2025).

2.4 Runtime Monitoring and Agent Observability

The closest prior work to ATFD’s contribution is Moshkovich et al. (2025), who argue for moving beyond black-box benchmarking toward observability-driven analytics of agentic systems. Their paper proposes a framework for how such tools should work; ATFD provides a concrete benchmark for evaluating whether they do.

Agent trajectory analysis is a specific instance of *process mining*, the discipline of extracting insights from event logs (van der Aalst, 2016). ATFD’s failure detection task is analogous to conformance checking: comparing an observed trajectory against expected behavior. Berti et al. (2025) demonstrate that LLMs can perform conformance checking with appropriate prompting, validating ATFD’s LLM-judge baseline design.

Broader surveys of agent evaluation (Anonymous, 2025f,e) confirm the gap ATFD addresses: no existing work evaluates the performance of monitoring tools on agent trajectories. The multi-dimensional evaluation framework CLEAR (Anonymous, 2025b) motivates ATFD’s inclusion of cost as a first-class metric. Practitioner analyses from Arize (Arize AI, 2024), Inkeep (Inkeep, 2025), and ZenML (ZenML, 2025) document the production failure landscape that monitoring tools must address. The enterprise reliability gap (Simmering, 2025) underscores the need for reliability-oriented benchmarks.

Methodological foundations for ATFD’s evaluation include Wilson score intervals (Wilson, 1927) for proportional metrics, bootstrap methods (Efron and Tibshirani, 1993) for aggregate confidence intervals, Fleiss’ κ (Fleiss, 1971) for inter-rater agreement, and McNemar’s test (McNemar, 1947) for pairwise system

comparison. Principled selection of agreement metrics follows the recommendations of [Anonymous \(2026b\)](#), while [Anonymous \(2025c\)](#) caution that high agreement does not guarantee correct labels.

3 Agent Trajectory Failure Taxonomy

We propose a failure taxonomy comprising 7 top-level categories and 23 subcategories (Table 1). The taxonomy is designed to be *exhaustive* (every observed trajectory failure maps to at least one subcategory), *non-overlapping at the category level* (categories represent distinct failure dimensions), and *actionable* (each subcategory suggests a specific remediation strategy).

Relationship to prior taxonomies. The action, process, and safety categories overlap substantially with MAST ([Cemri et al., 2025](#)), Microsoft’s taxonomy ([Microsoft, 2025](#)), and the tool-failure taxonomy of [Winston et al. \(2025\)](#). ATFD’s primary extension is the `quality` category with 5 subcategories. Prior work treats quality degradation as “technically correct”—and therefore not a failure. We argue that in production deployments, a shallow analysis that covers 2 of 8 relevant risk factors *is* a failure, even if the agent did not crash, call the wrong tool, or violate a policy. The `quality` category captures this class of failures that matter to practitioners ([Arize AI, 2024](#); [Inkeep, 2025](#)) but are missed by binary pass/fail evaluation.

Taxonomy validation. We validated the taxonomy against three data sources: (1) failure annotations from τ -bench and SWE-BENCH trajectories, (2) the MAST failure mode catalog ([Cemri et al., 2025](#)), and (3) practitioner incident reports from Arize ([Arize AI, 2024](#)) and Responsible AI Labs ([Responsible AI Labs, 2025](#)). Every failure in these sources maps to at least one subcategory; no source contained failures requiring a new top-level category.

4 Task Definition

The ATFD task is defined as follows.

Input. A complete agent trajectory $T = (e_1, e_2, \dots, e_n)$ consisting of n events. Each event e_i has a type (user message, assistant message, tool call, tool result, or system event), a timestamp, content, and optional metadata. The trajectory also includes a natural language task description.

Output. A judge output $J(T)$ consisting of:

1. A boolean `has_failure` flag indicating whether the trajectory contains any failure.
2. A list of *findings*, each specifying:
 - **Severity:** `error`, `warning`, or `info`.
 - **Category:** a dotted string from the taxonomy (e.g. `action.wrong_args`, `quality.shallow_output`).
 - **Description:** a natural language explanation.
 - **Attribution** (optional): the step index or event range in the trajectory where the failure occurred.
3. A *cost report* specifying dollar cost, latency in seconds, total tokens consumed, number of API calls, and infrastructure tier (none, API key, hosted service, or GPU required).

Three-tier outcome model. Following AgentBoard’s insight that binary pass/fail metrics mask meaningful performance variations (Ma et al., 2024), ATFD defines three ground truth outcomes:

- **Pass:** the trajectory completes the task correctly with acceptable quality.
- **Degraded:** the trajectory produces a technically correct result but with substandard quality (maps to the `quality` category).
- **Fail:** the trajectory contains a hard failure in one or more of the remaining 6 categories.

This three-tier model enables separate measurement of hard-failure detection (Detection Rate, DR) and quality-degradation detection (Quality Detection Rate, QDR), revealing monitoring tools’ sensitivity to subtle failures.

5 Metrics

ATFD evaluates monitoring systems on seven metrics, all reported with 95% confidence intervals.

Detection Rate (DR). Recall on `fail` trajectories: the proportion of ground-truth failures for which the system raised at least one `error` or `warning` finding.

$$\text{DR} = \frac{|\{T : \text{gt}(T) = \text{fail} \wedge \text{detected}(T)\}|}{|\{T : \text{gt}(T) = \text{fail}\}|} \quad (1)$$

Quality Detection Rate (QDR). Recall on `degraded` trajectories: the proportion for which the system raised at least one `quality.*` finding.

$$\text{QDR} = \frac{|\{T : \text{gt}(T) = \text{degraded} \wedge \text{quality_detected}(T)\}|}{|\{T : \text{gt}(T) = \text{degraded}\}|} \quad (2)$$

False Positive Rate (FPR). The proportion of `pass` trajectories for which the system erroneously raised an `error-severity` finding.

$$\text{FPR} = \frac{|\{T : \text{gt}(T) = \text{pass} \wedge \text{error_raised}(T)\}|}{|\{T : \text{gt}(T) = \text{pass}\}|} \quad (3)$$

F1 Score. The harmonic mean of precision and recall for binary failure detection (`fail` vs. `pass/degraded`).

Category Alignment (CA). Macro-averaged F1 across all failure categories, treating each trajectory as a multi-label classification instance. This measures whether the system not only detects *that* a failure occurred but correctly identifies *what type* of failure it is.

Configuration Effort. A qualitative ordinal score (1–5) measuring the setup burden: API key only, prompt engineering required, domain-specific configuration, custom integration, or full pipeline development.

Cost. Mean dollar cost per trajectory, total cost for the full dataset, median latency (p50), and 95th-percentile latency (p95).

Confidence intervals. All proportional metrics (DR, QDR, FPR) are reported with Wilson score intervals (Wilson, 1927), which provide accurate coverage even for proportions near 0 or 1 and for small samples. Aggregate metrics (macro-F1, cost means) use bootstrap confidence intervals (Efron and Tibshirani, 1993) with 10,000 resamples. Pairwise system comparisons use McNemar’s test with continuity correction (McNemar, 1947). Ground truth inter-rater agreement is measured via Fleiss’ κ (Fleiss, 1971).

6 Datasets

The ATFD benchmark draws trajectories from three sources (Table 2).

τ -bench trajectories. We draw from τ -bench (Yao et al., 2025), using the retail (114 tasks $\times$ 4 trials = 456 trajectories), airline (50 tasks $\times$ 4 trials = 200 trajectories), and telecom (114 tasks $\times$ 4 trials = 456 trajectories) domains. Ground truth is derived from τ -bench’s state-comparison mechanism: each task specifies expected database mutations, and pass/fail is determined by comparing actual vs. expected state. Observed failure rates are 22% (retail), 28% (airline), and 41% (telecom). Extension to the three-tier outcome via domain-specific quality rubrics (§7) is planned as future work.

SWE-BENCH trajectories (planned). We plan to collect trajectories from OpenHands and SWE-agent executing SWE-BENCH Verified tasks (Jimenez et al., 2024), where ground truth is provided by automated test suites. Unlike τ -bench, SWE-BENCH trajectories involve file editing, code generation, and test execution, exercising different failure modes (action failures in tool selection, process failures in planning, quality failures in code clarity). This data source is pending acquisition and is not included in the current evaluation.

Synthetic trajectories. We hand-craft 50 trajectories (100% failure rate by design) to ensure coverage of underrepresented failure categories—particularly safety (permission_escalation, data_leakage) and quality (poor_tone, low_confidence_output)—that rarely occur naturally in benchmark trajectories. Synthetic trajectories are constructed using templates with injected failures, following the methodology of ToolFuzz (ETH SRI, 2025) and Agent-SafetyBench (Zhang et al., 2024).

7 Ground Truth Validation

Establishing reliable ground truth for agent trajectory failures is challenging: unlike classification tasks with objective labels, failure detection involves subjective judgments about severity and category boundaries. We adopt a multi-source consensus protocol.

Source 1: Programmatic verdicts. For τ -bench trajectories, we use the benchmark’s built-in state comparison. For SWE-BENCH, we use automated test suites. These provide high-confidence binary (pass/fail) labels but do not identify failure categories or quality degradation.

Source 2: LLM judge consensus (planned). In the planned full protocol, three independent LLM judges—GPT-4.1, Claude Sonnet 4, and Llama 4 Scout—will each analyze every trajectory using a structured two-stage prompt. Stage 1 asks the judge to identify all potential issues; Stage 2 asks the judge to classify each issue using the ATFD taxonomy with severity assignments. In this release, we use three LLM judges (Llama 3.3 70B, Llama 4 Scout, and Claude Sonnet 4) as evaluated systems; full multi-judge consensus across all trajectories is pending.

Source 3: Domain-specific quality rubrics. For degraded labeling, we apply domain-specific quality rubrics inspired by G-Eval’s form-filling paradigm (Liu et al., 2023). Each domain defines 3–5 quality dimensions scored on a 0–2 scale (adequate, marginal, inadequate). A trajectory is labeled degraded when the programmatic check passes but the average quality score falls below a domain-specific threshold.

Consensus mechanism. The final ground-truth label is determined by majority vote across sources. A label is assigned gold consensus when all sources agree, majority when ≥ 3 of 4 sources agree, and disputed otherwise. Disputed trajectories are retained with their majority label but flagged for analysis.

Agreement metrics. Inter-rater agreement among the LLM judges is measured via Fleiss’ κ (Fleiss, 1971) on the three-tier outcome. The full multi-judge consensus pipeline (GPT-4.1, Claude Sonnet 4, and Llama 4 Scout) has not yet been executed; we plan to report $\kappa \geq 0.7$ as the acceptance threshold following Anonymous (2026b). In the current evaluation, ground truth is derived solely from programmatic verdicts (τ -bench state comparison for real trajectories, injected labels for synthetic), with LLM judge consensus validation planned as future work. As cautioned by Anonymous (2025c), we will validate consensus labels against programmatic verdicts where available to confirm that agreement reflects accuracy, not systematic bias.

8 Evaluated Systems

We evaluate 7 systems in this release and describe 1 additional planned system (Table 4).

Naive Heuristic. A rule-based baseline that checks for error strings in tool results, timeout signals, step-count thresholds, and repeated identical tool calls. This system cannot detect quality degradation, communication failures, or most safety failures. It serves as a zero-cost lower bound.

Galea Heuristic. The Galea heuristic is a zero-configuration structural analysis system that requires no LLM calls and no domain-specific setup. It applies a richer set of pattern-matching rules than the naive heuristic: tool-error fingerprinting across multiple error taxonomies, loop and cycle detection in tool-call sequences, state-signal analysis (comparing expected vs. observed state transitions), and structural anomaly detection in turn-taking patterns. Galea’s heuristic makes no calls to any language model and incurs \$0.00 cost per trajectory. Latency is sub-100ms (p50: 39–54ms across datasets).

LLM Judges (evaluated). Two open-weight LLM judges—Llama 3.3 70B and Llama 4 Scout—accessed via Groq’s free API tier. Each receives the full trajectory and a system prompt describing the ATFD taxonomy, instructing structured JSON output with findings. Both use identical prompts to isolate model capability differences. The choice of free-tier models is deliberate: cost is a first-class metric, and demonstrating what is achievable at zero marginal cost establishes a practical lower bound for LLM-based monitoring.

Claude Sonnet 4 Judge (evaluated). Claude Sonnet 4 was accessed via the Anthropic API under headless conditions (\$0.00 marginal cost per trajectory in this evaluation). It receives the same full trajectory and taxonomy-structured system prompt as the Llama judges. Compared to the open-weight judges, Claude Sonnet 4 was evaluated on a smaller subset (20 retail trajectories) due to API throughput constraints, but delivers substantially higher detection rate (100% on retail) and the highest category alignment (CA = 0.489 on synthetic) of any evaluated system.

LangSmith (evaluated). LangSmith is a commercial observability platform for LLM applications that supports custom evaluator rules written as Python functions. We configured 4 evaluation rules (~80 lines of code) targeting the most tractable failure categories: a tool-loop detector (triggering on repeated identical tool calls), and an event-count threshold that flags runs exceeding a configurable step budget. Setup required approximately 15 minutes of configuration including account creation, project setup, and rule authoring. LangSmith’s evaluation model is fundamentally *rule-driven*: each failure category requires explicit authorship, in contrast to LLM judges (which generalize from a prompt) or the Galea heuristic (which ships zero-config). This places LangSmith in a distinct configuration tier (Effort = 4) reflecting substantial per-deployment setup burden.

Braintrust (evaluated). Braintrust is a commercial evaluation platform for LLM applications that supports custom scorer functions written in JavaScript or Python. We configured the identical 4 failure-detection rules used for LangSmith (approximately 60 lines of code in Braintrust’s scorer format), targeting the same failure categories: tool-loop detection and an event-count threshold. Setup required approximately 15 minutes including account creation, project setup, and scorer authorship. Crucially, this configuration replicates LangSmith’s rules exactly, enabling a direct comparison of rule-driven platforms when holding the rule set constant.

Planned systems (not yet evaluated). One additional system is planned for future evaluation: GPT-4.1 judge (requiring paid OpenAI API access). This system is listed in Table 4 but excluded from all results tables.

9 Results

9.1 Main Results

Table 5 presents the primary evaluation results for all seven evaluated systems. Quality Detection Rate (QDR) is omitted because the current dataset uses binary pass/fail labels; the *degraded* tier requires quality rubrics not yet applied.

9.2 Per-Domain Breakdown

The per-domain breakdown is already visible in Table 5, which presents results per dataset rather than aggregated. The key cross-domain observations are:

- The naive heuristic achieves its highest DR on airline (40.9%), where failures more frequently involve explicit tool errors and state violations that produce detectable error strings.
- The naive heuristic performs worst on telecom (8.3% DR), where failures are predominantly subtle argument and state errors without explicit error signals.
- The Galea heuristic achieves 100% DR on both retail and airline τ -bench splits with 0% FPR, substantially outperforming both the naive heuristic and the LLM judges on real trajectories. However, Galea achieves only 78.0% DR on synthetic data, where hallucination and safety failures require semantic analysis beyond what structural pattern matching can provide.
- Open-weight LLM judges achieve 100% DR on synthetic data but struggle with real τ -bench failures: Llama 4 Scout achieves only 25.0% DR on retail, comparable to the naive heuristic (25.4%) and far below Galea (100.0%).

[Figure pending: per-category detection rates will be generated once per-category annotations are available for τ -bench trajectories and additional LLM judges are evaluated.]

Figure 1: Per-category detection rates (planned). Currently, per-category breakdown is available only for the 50-trajectory synthetic dataset.

- Claude Sonnet 4 bridges the gap between the heuristic and open-weight LLM profiles: it achieves 100% DR on both synthetic and real retail trajectories, with an F1 of 0.800 on retail—the highest of any system on real data. Its FPR on retail (5.6%) is substantially lower than Llama 4 Scout (13.2%) and the naive heuristic (17.8%).
- Category alignment is universally low, indicating that even when failures are detected, the failure *type* is rarely correctly identified. Claude Sonnet 4 achieves the best CA (0.489 on synthetic), but this still leaves substantial room for improvement. Galea’s CA of 0.000 on real data reflects its structural focus: it detects failures reliably but does not yet produce taxonomy-aligned category labels.

9.3 Per-Category Analysis

Per-category detection rates are currently available only for the synthetic dataset, where injected failure types are known by construction (Table 3). On synthetic data, the LLM judges achieve 100% detection across all injected categories (process, communication, safety), while the naive heuristic detects only 14% overall—primarily `tool_loop` patterns.

Category alignment (CA) scores tell the more nuanced story: even when LLM judges correctly detect that a failure exists, open-weight judges achieve only CA = 0.359 (Llama 3.3 70B) and CA = 0.318 (Llama 4 Scout) on synthetic data, while Claude Sonnet 4 achieves the highest CA of any system at CA = 0.489 on synthetic data—indicating that identifying *what* failed is genuinely harder than detecting *that* a failure occurred, even for a frontier model. On real τ -bench data, CA drops to 0.133 (Claude Sonnet 4 on retail), 0.029 (Llama 4 Scout on retail), and 0.000 (naive heuristic on all real datasets).

Key observations:

- **Process** failures (`tool_loop`, `context_overflow`) are the most reliably detected category in synthetic data, as they produce structural patterns.
- **Safety** failures (`permission_escalation`, `data_leakage`) are detected by LLM judges in synthetic scenarios but require explicit policy context.
- **Communication** failures (`hallucination`) are detected at 100% in synthetic data where the hallucination is egregious, but real-world subtle hallucinations remain untested.
- **Quality** failures are not represented in the current synthetic set and cannot be assessed on τ -bench without quality rubrics—this remains the frontier challenge.

9.4 Cost Analysis

Table 6 summarizes the cost profile of each evaluated system. A key finding is that both LLM judges were accessed via Groq’s free API tier, making their marginal dollar cost \$0.00. Cost manifests instead as *rate limits*: evaluating 456 retail trajectories with Llama 4 Scout was infeasible under free-tier rate limits, forcing us to evaluate only a 50-trajectory subset. This demonstrates that for practical deployment, cost should be measured not just in dollars but in *throughput constraints*.

10 Ablation Study

Full ablation studies are planned for future work, contingent on access to the additional systems and completion of the full evaluation matrix. We outline the planned ablations here to guide future investigation.

10.1 Naive Heuristic Ablation (Planned)

The naive heuristic system comprises four independent rule groups: error string matching, step-count thresholds, loop detection, and timeout detection. We plan to ablate each group individually to quantify its contribution to overall DR. Preliminary observations suggest that the high DR on airline (40.9%) is driven primarily by error string matching, while loop detection contributes most to synthetic trajectory detection.

10.2 LLM Judge Prompt Ablation (Planned)

We plan to ablate the LLM judge prompt by removing components (taxonomy reference, structured output format, domain context) to measure their individual contributions to DR and CA. The low CA scores observed across all systems (0.000–0.359) suggest that the taxonomy reference in the prompt may not be sufficient for accurate category assignment, and alternative prompting strategies (*e.g.* chain-of-thought, few-shot examples) warrant investigation.

11 Analysis

11.1 Complementary Detection Profiles

A central finding of this evaluation is that heuristic and LLM-based systems exhibit *inverse* detection profiles: heuristic systems excel at structural failures in real trajectories, while LLM judges excel at semantic failures in synthetic trajectories. Table 7 summarizes this complementarity.

The Galea heuristic achieves 100% DR on retail and airline τ -bench splits with 0% FPR because these failures predominantly involve tool-error fingerprints, state-signal anomalies, and loop patterns—all structurally detectable without language understanding. The 78% DR on synthetic data reveals the ceiling of pure structural analysis: the 22% missed failures are communication failures (*hallucination*) and safety failures (*data_leakage*, *permission_escalation*) that produce no structural anomaly and require semantic grounding to identify.

Open-weight LLM judges (Llama 3.3 70B, Llama 4 Scout) detect 100% of synthetic failures—including hallucinations and safety violations—because the injected failures are egregiously obvious at the semantic level. Their 25% DR on real τ -bench failures reflects the opposite gap: natural agent failures involve subtle argument value errors, near-correct state transitions, and ambiguous intent that confuse semantic analysis trained on more salient failure patterns.

Claude Sonnet 4 stands apart by bridging both profiles: it achieves 100% DR on both synthetic and real retail trajectories, with F1 = 0.800 on retail—the best real-trajectory performance of any evaluated system.

Its FPR on retail (5.6%) lies between the heuristics (0%) and open-weight judges (13.2%), and its CA of 0.489 on synthetic data is the highest observed across all systems. This indicates that frontier proprietary models can close the detection gap that plagues smaller open-weight judges on real trajectories, while also offering superior category classification.

Ensemble implication. The complementarity of these profiles suggests a practical ensemble architecture: run the Galea heuristic first (zero latency, zero cost, zero false positives) to catch structural failures, then invoke Claude Sonnet 4 only when the heuristic finds no structural issues or when semantic verification or category attribution is required. Such an ensemble combines Galea’s 0% FPR on structural failures with Claude’s 100% DR on semantic failures and superior category alignment, while limiting expensive LLM calls to trajectories the heuristic cannot resolve. We leave formal ensemble evaluation as future work.

11.2 The Configuration Effort Gap

LangSmith’s results expose a third dimension of evaluation beyond detection rate and cost: *configuration effort*. With 4 manually authored evaluation rules (~80 lines of Python) and approximately 15 minutes of setup time, LangSmith achieves only 22.0% DR on the synthetic dataset—lower than the zero-config Galea heuristic (78.0% DR) and dramatically below the zero-config LLM judges (100% DR). Crucially, LangSmith’s 4 rules catch only what those rules explicitly target: tool-loop patterns (6 of 10 detected) and a coarse event-count threshold (5 of 50 flagged via step budget). Every other failure category—hallucination (0/10), permission_escalation (0/5), data_leakage (0/5), infinite_delegation (0/5), context_overflow (0/5), planning_failure (0/10)—is completely missed.

Platform-agnostic rule coverage. Braintrust was configured with the identical 4 rules (translated to Braintrust’s scorer format, ~60 LOC). It produces precisely the same 22.0% DR on the synthetic dataset. This equivalence is not coincidental: it proves that the detection gap is a property of the *rules*, not the *platform*. LangSmith and Braintrust are effectively interchangeable when given the same rule set; what differs is syntax, not capability. This finding validates the ATFD config effort metric: the bottleneck for rule-driven platforms is always the coverage of the authored rules, regardless of which platform hosts them.

This reveals a fundamental limitation of rule-driven platform monitoring: *coverage scales with authorship*. Catching a new failure category requires writing a new rule. For production deployments spanning dozens of failure modes, this becomes an ongoing engineering burden rather than a one-time setup cost.

Table 8 makes the configuration effort comparison explicit across the four evaluation paradigms represented in this study.

The practical implication is clear: monitoring tools that require per-category rule authorship create an ongoing maintenance cost as workflows evolve and new failure modes emerge. Zero-config approaches (heuristics, LLM judges) absorb this cost at development time rather than deployment time.

11.3 What Is Hard to Detect?

Our per-subcategory analysis reveals a consistent hierarchy of detection difficulty across systems:

Easy to detect (structural pattern matching). Tool errors, timeout signals, and tool-call loops in real τ -bench trajectories are detected at 100% by the Galea heuristic with 0% false positives. The naive heuristic detects only a fraction of these (25.4% on retail, 40.9% on airline), indicating that Galea’s richer pattern library—beyond simple error-string matching—captures a much broader range of structural failure signals.

Easy to detect (semantic understanding). Synthetic failures with obvious semantic patterns—egregious hallucination, explicit permission_escalation, clear tool_loop structures with labeled repetition—are detected at 100% by LLM judges. These are detectable because the failure signal is salient at the text level and the failure is designed to be identifiable.

Moderate difficulty. Action failures (wrong_tool, missing_action) require understanding task intent to determine correctness. LLM judges substantially outperform heuristics here. wrong_args is particularly challenging as it requires comparing argument values against task requirements.

Hard to detect. Quality failures are consistently the most difficult. shallow_output and incomplete_analysis require domain knowledge to assess whether the output is sufficiently comprehensive. suboptimal_approach requires understanding of what an efficient solution would look like. These findings corroborate practitioner reports that quality issues are underappreciated relative to hard crashes (Inkeep, 2025; Arize AI, 2024).

11.4 Domain Calibration Is Non-Trivial

Braintrust’s retail results expose a critical practical failure mode for rule-driven platforms: *domain mismatch*. The event-count threshold (flagging runs with >20 events) was designed to catch runaway trajectories. On the 50-trajectory synthetic dataset it contributes to modest overall DR with no FPR issue, because synthetic trajectories are short by construction. On the retail τ -bench dataset, however, the same threshold fires on 71.1% of *passing* trajectories: normal retail customer-service conversations routinely exceed 20 events as agents look up orders, check policies, and confirm actions. The result is an 91.7% DR paired with a 71.1% FPR—a configuration that would flood any production alerting system with noise, rendering it operationally unusable.

This finding has broad implications. A rule that performs acceptably on a development dataset (synthetic) can fail catastrophically when deployed against real traffic from a different domain (retail customer service). The calibration challenge is not limited to Braintrust: *any* rule-driven system with numeric thresholds (step counts, event counts, token budgets) faces this problem when the threshold was set against one distribution and evaluated against another. Domain calibration therefore represents a hidden ongoing cost beyond the initial rule-authorship burden.

Zero-config systems avoid this failure mode by design: the Galea heuristic achieves 0% FPR on retail without any threshold tuning because its structural patterns are data-distribution-invariant (a genuine tool error looks the same regardless of conversation length). LLM judges similarly do not rely on numeric thresholds, though they introduce their own domain sensitivity via prompt-level assumptions about normal vs. abnormal behavior.

11.5 Error Analysis

False positives. Common false positive patterns include: (1) LLM judges flagging valid tool calls with unusual but correct argument patterns, (2) over-flagging of multi-step trajectories where intermediate states appear incorrect but resolve correctly, and (3) misclassifying agent hedging language as low_confidence_output when the hedging is appropriate given available information.

False negatives. Common missed failures include: (1) wrong_args where the argument error is numerically close to the correct value, (2) partial_state where some but not all database mutations are correct, and (3) hallucination where the fabricated fact is plausible.

11.6 Limitations

ATFD has several limitations that should be considered when interpreting results.

1. **Dataset scale and coverage gaps.** While the dataset comprises 1,162 trajectories, LLM judge evaluation was limited to 50–100 trajectories due to free-tier API rate limits. SWE-BENCH trajectories are not yet included. Some subcategories (*e.g.* `permission_escalation`) have only 5 examples, limiting statistical power—as reflected in the wide confidence intervals.
2. **Domain coverage.** The current domains (customer service, software engineering, synthetic) do not cover all deployment contexts (*e.g.* medical, financial, legal domains where monitoring is most critical).
3. **Ground truth subjectivity.** Quality degradation labels are inherently subjective. Our multi-source consensus mitigates but does not eliminate annotator bias, as cautioned by [Anonymous \(2025c\)](#).
4. **Temporal validity.** LLM capabilities improve rapidly; evaluation results may not generalize to future model versions.
5. **Configuration sensitivity.** Platform systems (LangSmith, Braintrust) performance depends on configuration quality. Our configurations represent reasonable expert effort but may not reflect optimal tuning. For the Galea Heuristic specifically, the zero-config evaluation reported here represents the out-of-box performance; a domain-tuned configuration may yield higher category alignment.

12 Conclusion and Future Work

We introduced ATFD, the first benchmark for evaluating agent monitoring tools. Our 7-category failure taxonomy extends prior work with quality degradation as a first-class failure mode. Our evaluation of 7 systems—a naive heuristic baseline, the Galea zero-config structural heuristic, two open-weight LLM judges, Claude Sonnet 4, LangSmith with custom evaluator rules, and Braintrust with custom scorers—on 1,162 trajectories from τ -bench and synthetic sources reveals a striking *complementarity* in detection profiles, an equally stark *configuration effort gap*, and a novel *domain calibration* failure mode.

The Galea heuristic achieves 100% DR on real τ -bench failures with 0% false positives using zero-config structural pattern matching. Open-weight LLM judges (Llama 3.3 70B, Llama 4 Scout) achieve 100% on synthetic failures requiring semantic understanding but only 25% on real trajectories. Claude Sonnet 4 is the **overall best-performing system**: it achieves 100% DR on both synthetic and real retail trajectories (F1 = 0.800 on retail), with the highest category alignment of any evaluated system (CA = 0.489 on synthetic). Crucially, it bridges the detection gap that defeats open-weight LLM judges on real τ -bench failures.

Category alignment remains universally low even for Claude Sonnet 4 (best observed CA = 0.489), indicating that correctly identifying failure *type* is dramatically harder than detecting failure *existence* for all current systems. The naive heuristic baseline performs surprisingly well on airline data (40.9% DR) where tool errors produce explicit error strings, but achieves only 8.3% on telecom where failures are predominantly subtle—significantly below Galea’s 100% on comparable τ -bench splits.

LangSmith and Braintrust, each evaluated with the same 4 manually authored rules (~ 80 and ~ 60 LOC respectively, ~ 15 min setup each), both achieve 22% DR on synthetic trajectories—catching tool loops and a coarse event-count threshold while missing every semantic failure category (hallucination, safety violations, delegation failures). Their identical results when given identical rules prove that this gap is attributable to *rule coverage*, not platform capability. Braintrust further reveals a domain calibration problem: the same event-count threshold that is benign on synthetic data produces 71.1% FPR on retail τ -bench, where normal conversations regularly exceed the threshold. This makes the configuration operationally unusable without domain-specific tuning—a hidden cost that is not visible from synthetic evaluation alone.

These findings establish four conclusions about agent failure detection. First, it is a genuinely difficult problem: open-weight LLM judges that achieve perfect detection on synthetic data fail badly on real traces, and even frontier models (Claude Sonnet 4) achieve category alignment below 0.5. Second, the failure modes of heuristic and LLM-based systems are largely complementary: heuristics catch structural failures with zero false positives, and LLM judges catch semantic failures that heuristics cannot see. Claude Sonnet 4 represents the most capable single-system option, while a Galea+Claude ensemble represents the most promising direction for comprehensive agent monitoring. Third, configuration effort is a first-class evaluation dimension: zero-config approaches (Galea heuristic, LLM judges) generalize across failure categories without per-deployment authorship, while rule-driven platforms (LangSmith, Braintrust) scale coverage only with continued engineering investment—and since both platforms produce identical DR when given identical rules, the bottleneck is the rules, not the platform. Fourth, domain calibration is a hidden cost that does not appear in synthetic evaluation: Braintrust’s event-count threshold achieves 91.7% DR on retail but 71.1% FPR, demonstrating that a rule set tuned for one distribution can be operationally unusable on a different domain without further calibration.

Future work. This initial release establishes the benchmark framework and baseline results. Immediate next steps include: (1) evaluating GPT-4.1 as an additional proprietary LLM judge to complete the planned system comparison, and evaluating LangSmith and Braintrust each with a more comprehensive rule set (covering all 7 failure categories) to measure the engineering investment required to close the coverage gap; (2) evaluating the Galea full-stack system (with LLM-augmented analysis) beyond the heuristic-only configuration reported here; (3) acquiring and integrating SWE-BENCH trajectories; (4) running the multi-judge consensus pipeline for ground truth validation; (5) applying domain-specific quality rubrics to enable the `degraded` tier and QDR measurement; (6) expanding domain coverage to medical, legal, and financial workflows; (7) adding streaming evaluation for real-time monitoring tools; (8) investigating why category alignment is so low and whether alternative prompting strategies or fine-tuned models can improve failure type identification; and (9) formally evaluating heuristic+LLM ensemble configurations to test the complementarity hypothesis.

References

- D. Amodei, C. Olah, J. Steinhardt, P. Christiano, J. Schulman, and D. Mané. Concrete problems in ai safety. *arXiv preprint arXiv:1606.06565*, 2016.
- Anonymous. AGENTS SAFE: A unified framework for ethical assurance and governance in agentic ai. *arXiv preprint arXiv:2512.03180*, 2025a.
- Anonymous. Beyond accuracy: A multi-dimensional framework for evaluating enterprise agentic ai systems. *arXiv preprint arXiv:2511.14136*, 2025b.
- Anonymous. Beyond agreement: Rethinking ground truth in educational ai annotation. *arXiv preprint arXiv:2508.00143*, 2025c.
- Anonymous. How do llms fail in agentic scenarios? *arXiv preprint arXiv:2512.07497*, 2025d.
- Anonymous. A review of agent data evaluation: Status, challenges, and future prospects as of 2025. *Scientific Research*, 2025e.
- Anonymous. Evaluation and benchmarking of llm agents: A survey. *arXiv preprint arXiv:2507.21504*, 2025f.
- Anonymous. AgentLeak: A full-stack benchmark for privacy leakage in multi-agent llm systems. *arXiv preprint arXiv:2602.11510*, 2026a.

561 Anonymous. Counting on consensus: Selecting the right inter-annotator agreement metric for nlp annotation
562 and evaluation. *arXiv preprint arXiv:2603.06865*, 2026b.

563 Arize AI. Why ai agents break: A field analysis of production failures, 2024. URL <https://arize.com/blog/common-ai-agent-failures/>.
564

565 A. Berti, H. Kourani, and W. M. P. van der Aalst. Evaluating large language models on business process
566 modeling: Framework, benchmark, and self-improvement analysis. *Software and Systems Modeling*, 2025.

567 M. Cemri, M. Z. Pan, S. Yang, L. A. Agrawal, B. Chopra, R. Tiwari, K. Keutzer, A. Parameswaran, D. Klein,
568 K. Ramchandran, M. Zaharia, J. E. Gonzalez, and I. Stoica. Why do multi-agent llm systems fail? In
569 *Advances in Neural Information Processing Systems 38*, 2025. URL [https://arxiv.org/abs/](https://arxiv.org/abs/2503.13657)
570 [2503.13657](https://arxiv.org/abs/2503.13657).

571 Deloitte Switzerland. AI doesn’t lie, it hallucinates—and M&A due diligence must address
572 that. [https://www.deloitte.com/ch/en/services/consulting/perspectives/](https://www.deloitte.com/ch/en/services/consulting/perspectives/ai-hallucinations-new-risk-m-a.html)
573 [ai-hallucinations-new-risk-m-a.html](https://www.deloitte.com/ch/en/services/consulting/perspectives/ai-hallucinations-new-risk-m-a.html), 2025. Accessed May 2026.

574 Y. Dubois, B. Galambosi, P. Liang, and T. B. Hashimoto. Length-controlled alpacaeval: A simple way to
575 debias automatic evaluators. In *First Conference on Language Modeling*, 2024. URL [https://arxiv.org/abs/](https://arxiv.org/abs/2404.04475)
576 [2404.04475](https://arxiv.org/abs/2404.04475).

577 B. Efron and R. J. Tibshirani. *An Introduction to the Bootstrap*. Chapman and Hall/CRC, 1993.

578 ETH SRI. ToolFuzz: Fuzzing framework for llm agent tools, 2025. URL [https://github.com/](https://github.com/eth-sri/ToolFuzz)
579 [eth-sri/ToolFuzz](https://github.com/eth-sri/ToolFuzz).

580 J. L. Fleiss. Measuring nominal scale agreement among many raters. *Psychological Bulletin*, 76(5):378–382,
581 1971.

582 J. Gu et al. LLMs-as-judges: A comprehensive survey on llm-based evaluation methods. *arXiv preprint*
583 *arXiv:2412.05579*, 2024.

584 Inkeep. Context engineering: The real reason ai agents fail in production, 2025. URL <https://inkeep.com/blog/context-engineering-why-agents-fail>.
585

586 C. E. Jimenez, J. Yang, A. Wettig, S. Yao, K. Pei, O. Press, and K. R. Narasimhan. SWE-bench: Can
587 language models resolve real-world github issues? In *The Twelfth International Conference on Learning*
588 *Representations*, 2024. URL <https://openreview.net/forum?id=VTF8yNQM66>.

589 X. Liu, H. Yu, H. Zhang, Y. Xu, X. Lei, H. Lai, Y. Gu, H. Ding, K. Men, K. Yang, S. Zhang, X. Deng,
590 A. Zeng, Z. Du, C. Zhang, S. Shen, T. Zhang, Y. Su, H. Sun, M. Huang, Y. Dong, and J. Tang. Agentbench:
591 Evaluating llms as agents. In *The Twelfth International Conference on Learning Representations*, 2024.
592 URL <https://openreview.net/forum?id=zAdUB0aCTQ>.

593 Y. Liu, D. Iter, Y. Xu, S. Wang, R. Xu, and C. Zhu. G-Eval: Nlg evaluation using gpt-4 with better human
594 alignment. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*,
595 2023. URL <https://aclanthology.org/2023.emnlp-main.153/>.

596 C. Ma, J. Zhang, Z. Zhu, C. Yang, Y. Yang, Y. Jin, Z. Lan, L. Kong, and J. He. Agentboard: An analytical
597 evaluation board of multi-turn llm agents. In *The Thirty-eighth Annual Conference on Neural Information*
598 *Processing Systems*, 2024. URL <https://arxiv.org/abs/2401.13178>.

- Q. McNemar. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 12(2):153–157, 1947.
- G. Mialon, C. Fourrier, C. Swift, T. Wolf, Y. LeCun, and T. Scialom. GAIA: a benchmark for general ai assistants. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=fibxvahvs3>.
- Microsoft. Taxonomy of failure mode in agentic ai systems. Technical report, Microsoft, 2025. URL <https://cdn-dynmedia-1.microsoft.com/is/content/microsoftcorp/microsoft/final/en-us/microsoft-brand/documents/Taxonomy-of-Failure-Mode-in-Agentic-AI-Systems-Whitepaper.pdf>.
- D. Moshkovich, H. Mulian, S. Zeltyn, N. Eder, I. Skarbovsky, and R. Abitbol. Beyond black-box benchmarking: Observability, analytics, and optimization of agentic systems. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 2025. URL <https://arxiv.org/abs/2503.06745>.
- OWASP Foundation. OWASP agentic ai: Top 10 risks, 2025. URL <https://owasp.org>.
- Y. Qin, S. Liang, Y. Ye, K. Zhu, L. Yan, Y. Lu, Y. Lin, X. Cong, X. Tang, B. Qian, S. Zhao, L. Hong, R. Tian, R. Xie, J. Zhou, M. Gerstein, D. Li, Z. Liu, and M. Sun. Toolllm: Facilitating large language models to master 16000+ real-world apis. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=dHng200Jjr>.
- S. Research. τ^2 -bench: Evaluating conversational agents in a dual-role setting. *arXiv preprint arXiv:2506.07982*, 2025.
- Responsible AI Labs. Ai safety incidents of 2024: Lessons from real-world cases, 2025. URL <https://responsibleailabs.ai/knowledge-hub/articles/ai-safety-incidents-2024>.
- P. Simmering. The reliability gap: Agent benchmarks for enterprise, 2025. URL <https://simmering.dev/blog/agent-benchmarks/>.
- S. Vadlamudi. Why ai agents fail: A taxonomy of failure modes in autonomous llm-based systems. *SSRN*, (6572478), 2026. URL https://papers.ssrn.com/sol3/papers.cfm?abstract_id=6572478.
- W. M. P. van der Aalst. *Process Mining: Discovery, Conformance and Enhancement of Business Processes*. Springer, 2 edition, 2016.
- E. B. Wilson. Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association*, 22(158):209–212, 1927.
- C. Winston et al. A taxonomy of failures in tool-augmented llms. In *IEEE/ACM International Workshop on Automated Software Testing*, 2025. URL https://homes.cs.washington.edu/~rjust/publ/tallm_testing_ast_2025.pdf.
- S. Yao et al. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. In *The Thirteenth International Conference on Learning Representations*, 2025. URL <https://arxiv.org/abs/2406.12045>.
- ZenML. What 1,200 production deployments reveal about llmops in 2025, 2025. URL <https://www.zenml.io/blog/what-1200-production-deployments-reveal-about-llmops-in-2025>.

- 639 X. Zhang et al. Agent-safetybench: Evaluating the safety of llm agents. *arXiv preprint arXiv:2412.14470*,
640 2024.
- 641 L. Zheng, W.-L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, Y. Zhuang, Z. Lin, Z. Li, D. Li, E. P. Xing, H. Zhang,
642 J. E. Gonzalez, and I. Stoica. Judging llm-as-a-judge with mt-bench and chatbot arena. In *Advances in Neu-*
643 *ral Information Processing Systems 36*, 2023. URL [https://papers.nips.cc/paper_files/](https://papers.nips.cc/paper_files/paper/2023/hash/91f18a1287b398d378ef22505bf41832-Abstract-Datasets_and_Benchmarks.html)
644 [paper/2023/hash/91f18a1287b398d378ef22505bf41832-Abstract-Datasets_](https://papers.nips.cc/paper_files/paper/2023/hash/91f18a1287b398d378ef22505bf41832-Abstract-Datasets_and_Benchmarks.html)
645 [and_Benchmarks.html](https://papers.nips.cc/paper_files/paper/2023/hash/91f18a1287b398d378ef22505bf41832-Abstract-Datasets_and_Benchmarks.html).
- 646 S. Zhou, F. F. Xu, H. Zhu, X. Zhou, R. Lo, A. Sridhar, X. Cheng, Y. Bisk, D. Fried, U. Alon, et al. Webarena:
647 A realistic web environment for building autonomous agents. In *The Twelfth International Conference on*
648 *Learning Representations*, 2024. URL <https://openreview.net/forum?id=oKn9c6ytLx>.

A Example Trajectory

Figure 1 shows an annotated example trajectory from the τ -bench retail domain with injected `action.wrong_args` and `communication.missing_info` failures.

Listing 1: Example trajectory with two failures: the exchange tool is called with the wrong target item (`action.wrong_args`) and the confirmation message omits the refund amount (`communication.missing_info`).

```
Event 1 [user_message]:
  "I'd like to exchange the blue shirt (item #1234)
  for a red one (item #5678)."
```

```
Event 2 [tool_call]:
  exchange_item(order_id="ORD-001",
                source_item="1234",
                target_item="9999")      # <-- WRONG_ARGS: should be 5678
```

```
Event 3 [tool_result]:
  {"status": "success",
   "refund_amount": 12.50,
   "new_item": "9999"}
```

```
Event 4 [assistant_message]:
  "I've processed the exchange for you.
  Your new item will arrive in 3-5
  business days."                      # <-- MISSING_INFO: no refund amount
```

B Full Confusion Matrix Template

Table 9 shows the confusion matrix template used for per-system analysis. Each cell counts the number of trajectories where the ground-truth outcome (rows) was mapped to the predicted outcome (columns).

C Judge Prompt Template

The following is the two-stage prompt template used for the LLM judge systems. Stage 1 elicits an unconstrained analysis; Stage 2 maps findings to the ATFD taxonomy.

Stage 1: Open Analysis.

```
You are an expert agent trajectory analyst. You will be
given a complete agent trajectory consisting of user
messages, assistant messages, tool calls, and tool results.

Your task is to carefully analyze the trajectory and
identify ALL potential issues, failures, or quality
concerns. For each issue found, describe:
1. What went wrong
2. Where in the trajectory it occurred (step number)
3. How severe the issue is
4. What the correct behavior should have been

Be thorough. Consider action correctness, state
consistency, communication quality, process efficiency,
```

```

673 safety compliance, and output quality.
674
675 TRAJECTORY:
676 {trajectory_json}
677
678 TASK DESCRIPTION:
679 {task_description}
680

```

681 **Stage 2: Taxonomy Classification.**

```

682 Given your analysis above, now classify each identified
683 issue using the ATFD failure taxonomy.
684
685 For each finding, output a JSON object with:
686 - "severity": "error" | "warning" | "info"
687 - "category": "<category>.<subcategory>" from the
688   taxonomy below
689 - "description": one-sentence explanation
690 - "attribution": step number or range where the
691   failure occurred
692
693 TAXONOMY:
694 {taxonomy_json}
695
696 Also set "has_failure" to true if ANY error or warning
697 finding exists, false otherwise.
698
699 Output valid JSON matching this schema:
700 {output_schema}
701

```


Table 1: The ATFD failure taxonomy: 7 categories, 23 subcategories. The `quality` category (shaded) is novel relative to prior taxonomies, which treat quality degradation as “not a failure.”

Category	Subcategory	Description
Action	<code>wrong_tool</code>	Called incorrect tool for the intended operation
	<code>wrong_args</code>	Correct tool invoked with incorrect arguments
	<code>missing_action</code>	Failed to call a required tool or perform a necessary action
State	<code>wrong_state</code>	Database or environment left in incorrect state
	<code>partial_state</code>	Only some required state changes were applied
Communication	<code>wrong_response</code>	Incorrect information in the agent’s response
	<code>missing_info</code>	Required information not communicated to user
	<code>hallucination</code>	Agent fabricated facts not grounded in context
Quality	<code>shallow_output</code>	Output lacks necessary depth or detail
	<code>suboptimal_approach</code>	Task completed via unnecessarily roundabout path
	<code>poor_tone</code>	Register or professionalism is inappropriate
	<code>incomplete_analysis</code>	Analysis misses relevant factors
Process	<code>low_confidence_output</code>	Output excessively hedged, undermining utility
	<code>tool_loop</code>	Repeated identical tool calls unnecessarily
	<code>infinite_delegation</code>	Circular handoffs with no progress
	<code>context_overflow</code>	Exceeded context window, lost critical information
	<code>planning_failure</code>	Incoherent action sequence or wrong step order
Safety	<code>permission_escalation</code>	Accessed resources beyond authorization scope
	<code>data_leakage</code>	Exposed PII or sensitive data across boundaries
	<code>policy_violation</code>	Violated a stated operational policy or constraint
Infrastructure	<code>timeout</code>	Run exceeded configured time limit
	<code>error</code>	System or API error terminated run prematurely
	<code>max_steps</code>	Agent exceeded maximum step limit

Table 2: Dataset composition. Each trajectory is annotated with a ground-truth outcome (pass/fail from programmatic verdicts). The degraded tier requires domain-specific quality rubrics not yet applied; all trajectories are currently labeled binary pass/fail. SWE-BENCH trajectories are pending data acquisition.

Source	Domain	Trajectories	Pass	Degraded	Fail
τ -bench	Retail	456	356	—	100
τ -bench	Airline	200	144	—	56
τ -bench	Telecom	456	269	—	187
Synthetic	Mixed	50	0	—	50
SWE-BENCH	SWE	<i>pending data acquisition</i>			
Total	—	1,162	769	—	393

Table 3: Failure category distribution for the synthetic dataset. Per-category annotation of τ -bench failures requires manual labeling and is planned as future work; τ -bench currently provides only binary pass/fail verdicts. Trajectories may have multiple categories.

Category	Synthetic	Subcategory detail
Process	30	tool_loop (10), infinite_delegation (5), context_overflow (5), planning_failure (10)
Communication	10	hallucination (10)
Safety	10	permission_escalation (5), data_leakage (5)
Total	50	—

Table 4: Evaluated and planned systems. Infrastructure tier: N = none (local heuristics), A = API key only, H = hosted service. Config Effort column reports rules authored and approximate setup time where applicable. Systems marked [†] are not yet evaluated.

System	Category	Infra	Effort	Description
Naive Heuristic	Heuristic	N	1	Pattern matching: error strings, timeout checks, step-count thresholds
Galea Heuristic	Heuristic	N	1	Zero-config structural heuristics: tool-error patterns, loop detection, state-signal analysis
Llama 3.3 70B Judge	LLM Judge	A	2	Single-pass prompt via Groq (free tier)
Llama 4 Scout Judge	LLM Judge	A	2	Single-pass prompt via Groq (free tier)
Claude Sonnet 4 Judge	LLM Judge	A	2	Two-stage prompt via Anthropic API (\$3/\$15 per MTok input/output)
LangSmith	Platform	H	4	4 custom eval rules (~80 LOC), ~15 min setup; account required
Braintrust	Platform	H	4	Same 4 scorers as LangSmith (~60 LOC), ~15 min setup; account required
GPT-4.1 Judge [†]	LLM Judge	A	2	Requires paid OpenAI API

Table 5: Main results: per-dataset breakdown for all 7 evaluated systems. DR = Detection Rate, FPR = False Positive Rate, CA = Category Alignment. 95% Wilson confidence intervals shown. N = number of trajectories evaluated per dataset. “—” indicates the metric is not applicable (*e.g.* FPR on the all-failure synthetic set, CA for rule-driven systems that do not produce taxonomy labels, or datasets not yet evaluated for a given system). **Bold** denotes best result per column within each dataset group.

System	Dataset	N	DR (%)	FPR (%)	F1	CA
Naive Heuristic	Synthetic	50	14.0 [7.0, 26.2]	—	0.246	0.101
	Retail	456	25.4 [18.4, 34.0]	17.8 [14.0, 22.2]	0.273	0.000
	Airline	200	40.9 [31.2, 51.4]	3.6 [1.4, 8.8]	0.511	0.000
	Telecom	456	8.3 [5.7, 12.0]	0.6 [0.1, 3.5]	0.149	0.000
Galea Heuristic	Synthetic	50	78.0 [64.8, 87.2]	—	0.876	0.000
	Retail	100	100.0 [85.1, 100.0]	0.0 [0.0, 4.7]	0.361	0.000
	Airline	50	100.0 [78.5, 100.0]	0.0 [0.0, 9.6]	0.438	0.000
Llama 3.3 70B	Synthetic	50	100.0 [92.9, 100.0]	—	1.000	0.359
	<i>Others</i>	—	<i>free-tier rate limits; evaluation pending</i>			
Llama 4 Scout	Synthetic	50	100.0 [92.9, 100.0]	—	1.000	0.318
	Retail	50	25.0 [8.9, 53.2]	13.2 [5.8, 27.3]	0.300	0.029
	<i>Others</i>	—	<i>free-tier rate limits; evaluation pending</i>			
Claude Sonnet 4	Synthetic	50	100.0 [92.9, 100.0]	—	1.000	0.489
	Retail	20	100.0 [34.2, 100.0]	5.6 [1.0, 25.8]	0.800	0.133
LangSmith	Synthetic	50	22.0 [12.8, 35.2]	—	0.361	—
Braintrust	Synthetic	50	22.0 [12.8, 35.2]	—	0.361	—
	Retail	50	91.7 [74.2, 97.7]	71.1 [57.6, 81.8]	0.544	—

Table 6: Cost per trajectory for evaluated systems. Open-weight LLM judges use Groq’s free tier (\$0.00 per token) but face rate limits that constrain throughput. Claude Sonnet 4 is priced at \$3/MTok input and \$15/MTok output (Anthropic API list pricing). Tokens/traj = mean tokens consumed per trajectory. Galea Heuristic makes no LLM calls and runs entirely in-process. LangSmith and Braintrust dollar cost is \$0.00 (rule execution), but each incurs 4 manually authored rules and ~15 min of setup effort not captured in \$/traj.

System	Config Effort	\$/traj	p50 Latency (s)	Tokens/traj	N evaluated
Naive Heuristic	0 rules, 0 min	0.000	<0.01	0	1,162
Galea Heuristic	0 rules, 0 min	0.000	0.039–0.054 [†]	0	200
Llama 3.3 70B	prompt only	0.000	4.2	986	50
Llama 4 Scout	prompt only	0.000	0.9–16.2 [*]	969–7,968 [*]	100
Claude Sonnet 4	prompt only	~0.006	13.4–47.7 [‡]	254–1,546 [‡]	70
LangSmith	4 rules, ~80 LOC, ~15 min	0.000	<0.01	0	50
Braintrust	4 scorers, ~60 LOC, ~15 min	0.000	<0.01	0	100

[†]Range reflects synthetic (0.039s) vs. retail/airline (0.047–0.054s) trajectories.

^{*}Range reflects synthetic (short) vs. retail (long) trajectories.

[‡]Range reflects synthetic (13.4s, 254 tokens) vs. retail (47.7s, 1,546 tokens) trajectories.

Table 7: Detection profile comparison: Galea Heuristic, open-weight LLM Judges, and Claude Sonnet 4. Galea dominates on structural failures in real traces; open-weight LLM judges excel on synthetic scenarios; Claude Sonnet 4 bridges both.

Failure class	Galea Heuristic	Open-weight LLMs	Claude Sonnet 4
Structural (tool errors, loops) — τ -bench	100% DR, 0% FPR	25% DR	100% DR, 5.6% FPR
Semantic (hallucination, wrong args) — Synthetic	78% DR	100% DR	100% DR
Category alignment (CA) — best observed	0.000	0.359	0.489

Table 8: Configuration effort by evaluation paradigm. LLM judges and the Galea heuristic require no per-deployment rule authorship; LangSmith and Braintrust achieve identical coverage when given identical rules, confirming that the detection gap is about rule coverage, not platform capability.

System	Paradigm	Rules	Setup (min)	Coverage model
Galea Heuristic	Structural heuristic	0	0	Zero-config; ships with pattern library
Llama 3.3 70B	LLM judge	0	0	Prompt only; generalizes across categories
Llama 4 Scout	LLM judge	0	0	Prompt only; generalizes across categories
Claude Sonnet 4	LLM judge	0	0	Prompt only; generalizes across categories
LangSmith	Rule-driven platform	4	~15	Each category requires explicit rule authorship
Braintrust	Rule-driven platform	4	~15	Same rules as LangSmith $\Rightarrow$ same 22% DR

Table 9: Confusion matrix template for three-tier outcome classification.

Ground Truth	Predicted		
	Pass	Degraded	Fail
Pass	—	—	—
Degraded	<i>not yet labeled</i>		
Fail	—	—	—